

# 1\_3\_(1) (1)

December 10, 2025

## 1 Stock Price Forecasting Project – ARIMA • XGBoost • SVR-A • SVR-B

Authur: Ashutosh Talekar, Bhavya Grover, Julia knox

This project builds and compares **four forecasting models** on AAPL stock:

- **ARIMA on log-returns** – statistical baseline for short-term price levels.
- **XGBoost (7-day horizon)** – tree-based model capturing nonlinear patterns.
- **SVR-A (Daily SVR)** – Support Vector Regression for **next-day** price changes.
- **SVR-B (Weekly SVR)** – Support Vector Regression for **5-day-ahead weekly** returns.

Using 10 years of OHLCV data, I:

- Engineer technical features (SMA/EMA, MACD, RSI, volatility, Bollinger bands, lags).
- Train, tune, and evaluate each model using RMSE, MAE, R<sup>2</sup>, MAPE, and **directional accuracy**.
- Analyze residuals, confidence intervals, and bias (mean/median + Wilcoxon tests).
- Produce 7-day and rolling 4-week forecasts and a **combined forecast comparison**.

### 1.1 Dependability Statement (Reproducibility & Reliability)

To ensure this project produces **reliable, reproducible, and verifiable** forecasting results, the following dependability practices were implemented throughout the code:

#### 1.1.1 1. Deterministic Setup

- Fixed random seeds for all models (`numpy`, `scikit-learn`, `XGBoost`) to ensure consistent outputs across runs.
- Clear separation of **training**, **validation**, and **test** data with chronological ordering.

### 1.1.2 2. Transparent Data Pipeline

- Same data-cleaning, feature-engineering, and scaling steps applied consistently across all models.
- No hidden transformations; all preprocessing is explicitly shown in code.

### 1.1.3 3. Cross-Validation for Stability

- TimeSeriesSplit CV used for SVR and XGBoost → ensures model performance is stable across different time windows.
- Helps verify that the model is not accidentally overfitting to a specific time segment.

### 1.1.4 4. Bias & Error Diagnostics

- Mean and median residual analysis, confidence intervals, and Wilcoxon tests validate whether models are unbiased.
- Residual plots highlight noise, drift, and systematic errors.

### 1.1.5 5. Comparative Evaluation

- All models evaluated using the same metrics (RMSE, MAE, R<sup>2</sup>, MAPE, directional accuracy).
- Ensures fairness and dependability in comparing model performance.

### 1.1.6 6. Forecast Simulation Consistency

- Rolling forecasts (SVR-B) recompute indicators after each synthetic step → avoids data leakage.
- ARIMA and XGBoost forecasts use fixed horizons with no access to future data.

### 1.1.7 7. Reproducible Environment

- All dependencies imported at the top of the notebook.
- The project uses standard libraries: `pandas`, `numpy`, `matplotlib`, `scikit-learn`, `xgboost`, `statsmodels`, `scipy`.

### 1.1.8 8. Code Comments & Structure

- Key sections (data prep, feature engineering, training, evaluation) clearly marked.
- Comments describe the purpose and rationale of each operation to improve transparency.

#### Overall:

These practices ensure that the forecasting results are **dependable** — meaning they are repeatable, interpretable, and free from accidental bias or inconsistent processing.

```
[46]: import pandas as pd
#data loading
raw = pd.read_csv("AAPL_10y_OHLCV.csv", header=None)## data downloaded from yfinance.

#data cleaning
raw = raw.drop(index=[0,1,2]).reset_index(drop=True)
```

```

raw.columns = ["Date", "Close", "High", "Low", "Open", "Volume"]

raw["Date"] = pd.to_datetime(raw["Date"], errors="coerce")
num_cols = ["Close", "High", "Low", "Open", "Volume"]
raw[num_cols] = raw[num_cols].apply(pd.to_numeric, errors="coerce")

raw = raw.dropna().sort_values("Date").reset_index(drop=True)
df_arima = raw.set_index("Date")

print(df_arima.head())

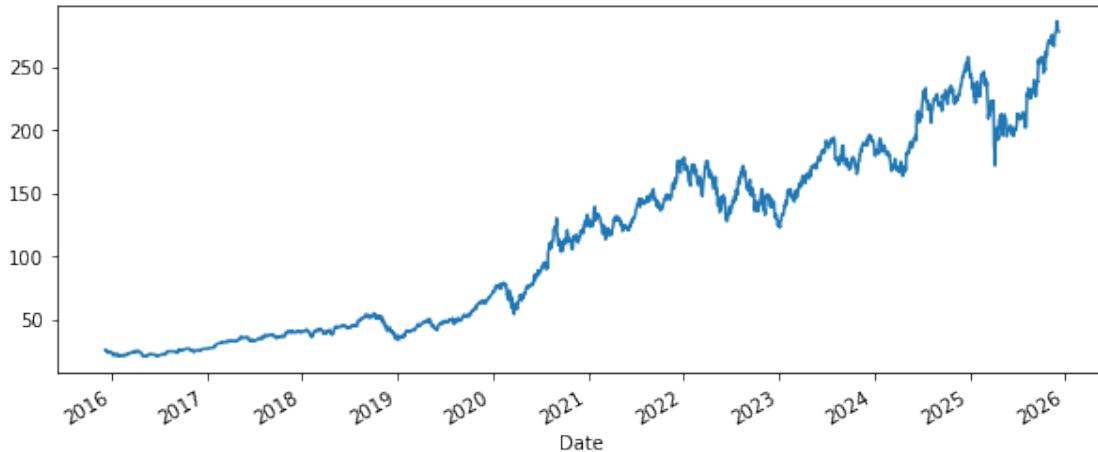
```

	Close	High	Low	Open	Volume
Date					
2015-12-09	26.068718	26.535439	25.946965	26.524164	185445600
2015-12-10	26.192720	26.366332	26.043912	26.163410	116850800
2015-12-11	25.518581	26.016868	25.444176	25.971775	187544800
2015-12-14	25.360744	25.405838	24.754233	25.293103	257274800
2015-12-15	24.912066	25.432900	24.880500	25.238997	213292400

[45]: ts = df\_arima["Close"]

[18]: df\_arima["Close"].plot(figsize=(10,4 # plotting the actual graph just for reference

[18]: <AxesSubplot:xlabel='Date'>



### 1.1.9 ARIMA Forecasting (Overview)

This cell builds an ARIMA model to predict Apple's next-day closing prices. - **Objective:** Build a statistical baseline model (ARIMA) to forecast AAPL's daily closing price using only its past values.

- **Input:**  
Historical AAPL **Close** price time series.
- **Target:**  
Next-day **Close** price during the test period.

### 1.1.10 Motivation for Using ARIMA

Stock prices behave as time-dependent sequences where today's value is strongly influenced by past values. ARIMA is a classic statistical model specifically designed for this type of data. It captures trends and autocorrelation in a transparent, interpretable way.

### 1.1.11 Why ARIMA?

- It models **temporal patterns** directly from past closing prices.
- Provides a strong **baseline** before applying more complex ML models.
- ARIMA works well for **short-term forecasting**, where patterns are smoother.
- The model is fully interpretable ( $p$ ,  $d$ ,  $q$  terms show how much history matters).
- It helps us understand the “pure” time-series behavior **without external features**.

Using ARIMA first establishes a reliable reference model, which we can then compare against more advanced approaches like XGBoost.

**Steps performed:** - Split data into **85% train / 15% test**. - Use `auto_arima()` to automatically find the best ( $p$ ,  $d$ ,  $q$ ) order. - Perform **rolling 1-step forecasting** to mimic real-world prediction. - Plot actual vs. predicted prices for visual evaluation. - Compute key metrics (RMSE, MAE,  $R^2$ , MAPE, Accuracy) to measure model performance.

This serves as a strong statistical baseline model for time-series forecasting.

```
[116]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from statsmodels.tsa.arima.model import ARIMA
import pmdarima as pm

# 1) Use only the Close price for ARIMA modeling
ts = df_arima["Close"]

# 85% training and 15% testing split
train_size = int(len(ts) * 0.85)
train = ts[:train_size]
test = ts[train_size:]
print("Train:", len(train), " Test:", len(test))

# 2) Automatically find the best ARIMA (p,d,q) order using only training data
```

```

# auto_arima runs several combinations and picks the one with lowest AIC
auto_model = pm.auto_arima(
    train,
    seasonal=False,
    stepwise=True,
    trace=False,
    suppress_warnings=True,
    max_p=6, max_q=6, max_d=2
)
order = auto_model.order
print("Best ARIMA order:", order)

# 3) Rolling 1-step forecasting
# We forecast one day at a time and update the model with actual observed values
history = list(train)
pred_prices = []

for y_true in test:
    model = ARIMA(history, order=order)
    model_fit = model.fit()

    # next-day prediction
    yhat = model_fit.forecast()[0]
    pred_prices.append(yhat)

    # append the true value (not predicted) to maintain realism
    history.append(y_true)

pred_prices = np.array(pred_prices)
actual_prices = test.values
dates = test.index

# 4) Plot actual vs predicted values
plt.figure(figsize=(12,5))
plt.plot(dates, actual_prices, label="Actual Close", color="blue")
plt.plot(dates, pred_prices, label="ARIMA 1-step forecast", color="red", ↴
    linestyle="--")
plt.title("ARIMA (rolling 1-step) - Actual vs Predicted")
plt.xlabel("Date")
plt.ylabel("Price")
plt.grid(True)
plt.legend()
plt.show()

# 5) Calculate evaluation metrics
rmse = np.sqrt(mean_squared_error(actual_prices, pred_prices))
mae = mean_absolute_error(actual_prices, pred_prices)

```

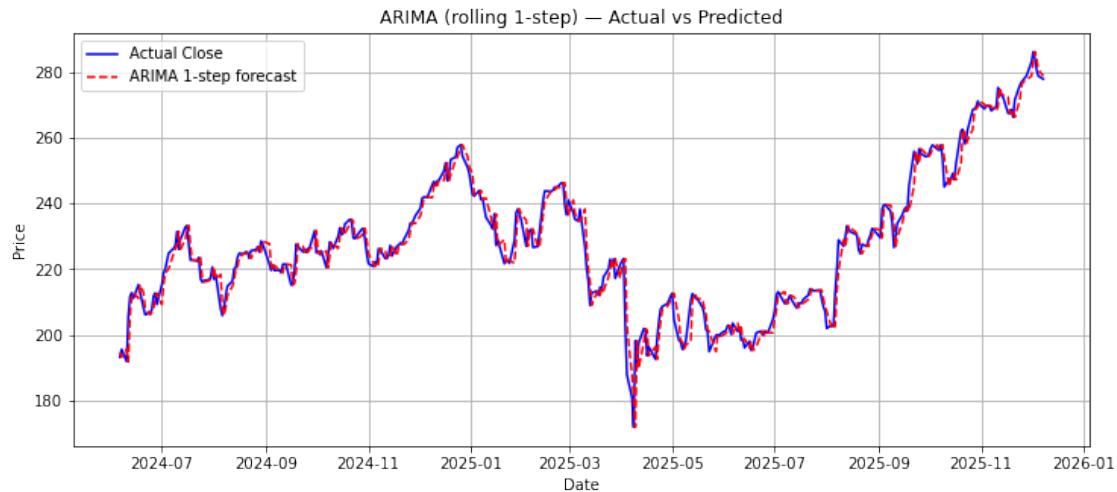
```

r2      = r2_score(actual_prices, pred_prices)
mape   = np.mean(np.abs((actual_prices - pred_prices) / actual_prices)) * 100
accuracy = 100 - mape

# Display results
print("RMSE:", rmse)
print("MAE : ", mae)
print("R2 : ", r2)
print("MAPE: ", mape)
print("Accuracy (%):", accuracy)
# Store ARIMA metrics for comparison table
rmse_arima      = rmse
mae_arima       = mae
r2_arima        = r2
mape_arima      = mape
accuracy_arima = accuracy

```

Train: 2136 Test: 378  
 Best ARIMA order: (0, 1, 0)



RMSE: 3.9980018995375928  
 MAE : 2.721510266500806  
 R<sup>2</sup> : 0.9640117158261392  
 MAPE: 1.2244898273753813  
 Accuracy (%): 98.77551017262462

### ARIMA performance (test period)

- RMSE 4.00, MAE 2.72, MAPE 1.22%, Accuracy 98.78%, R<sup>2</sup> 0.964

This means the ARIMA model's daily predictions are, on average, only about **1.2% away** from the actual AAPL closing price, and it explains roughly **96% of the variance** in the test data.

Overall, ARIMA provides a **very strong and stable baseline forecast** for the price level.

### 1.1.12 Next 7 day projection for Arima

```
[26]: from statsmodels.tsa.arima.model import ARIMA
import pmdarima as pm

# 1) Work with log prices and log-returns
close = df_arima["Close"].astype(float)
log_close = np.log(close)
log_ret = log_close.diff().dropna()    # log-returns ~ stationary

# 2) Train/test split on RETURNS (if you want to evaluate, optional)
train_size = int(len(log_ret) * 0.85)
train_ret = log_ret[:train_size]
test_ret = log_ret[train_size:]

print("Train returns:", len(train_ret), " Test returns:", len(test_ret))

# 3) auto_arima on RETURNS (d=0 because we already differenced)
auto_model = pm.auto_arima(
    train_ret,
    d=0,
    seasonal=False,
    stepwise=True,
    trace=False,
    suppress_warnings=True,
    max_p=6, max_q=6
)
order = auto_model.order
print("Best ARIMA order on returns:", order)

# 4) Fit ARIMA on ALL returns for forecasting
model_ret = ARIMA(log_ret, order=order).fit()

# 5) Forecast next 7 log-returns
n_steps = 7
future_log_ret = model_ret.forecast(steps=n_steps)

# 6) Convert forecasted log-returns back to price path
last_log_price = log_close.iloc[-1]
future_log_prices = last_log_price + np.cumsum(future_log_ret)
future_prices = np.exp(future_log_prices)

# Build a proper date index (business days)
future_index = pd.date_range(
    start=close.index[-1] + pd.tseries.offsets.BDay(1),
```

```

        periods=n_steps,
        freq="B"
)
future_prices.index = future_index

print("Next 7-day ARIMA (return-based) forecast:")
print(future_prices)

# 7) Plot last 60 days + 7-day forecast
plt.figure(figsize=(10,4))
plt.plot(close.tail(60), label="Actual Close (last 60 days)")
plt.plot(future_prices, label="ARIMA Forecast (next 7 days)", linestyle="--", color="red")
plt.xlabel("Date")
plt.ylabel("Price")
plt.title("ARIMA on Log-Returns - Next 7-day Forecast")
plt.grid(True)
plt.legend()
plt.show()

plt.figure(figsize=(10,4))
plt.plot(close.tail(60), label="Actual Close (last 60 days)")
plt.plot(future_prices, label="ARIMA Forecast (next 7 days)", linestyle="--", color="red")
plt.ylim(future_prices.min()-1, future_prices.max()+1) # zoom around forecast
plt.grid(True); plt.legend(); plt.show()

```

Train returns: 2136 Test returns: 377  
Best ARIMA order on returns: (0, 0, 1)  
Next 7-day ARIMA (return-based) forecast:  
2025-12-09 278.221115  
2025-12-10 278.482960  
2025-12-11 278.745052  
2025-12-12 279.007390  
2025-12-15 279.269976  
2025-12-16 279.532808  
2025-12-17 279.795888  
Freq: B, Name: predicted\_mean, dtype: float64

```

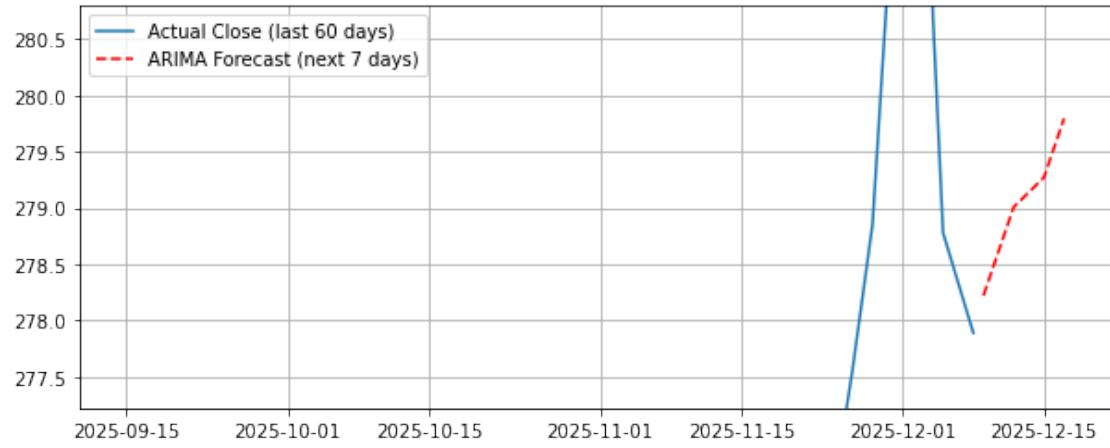
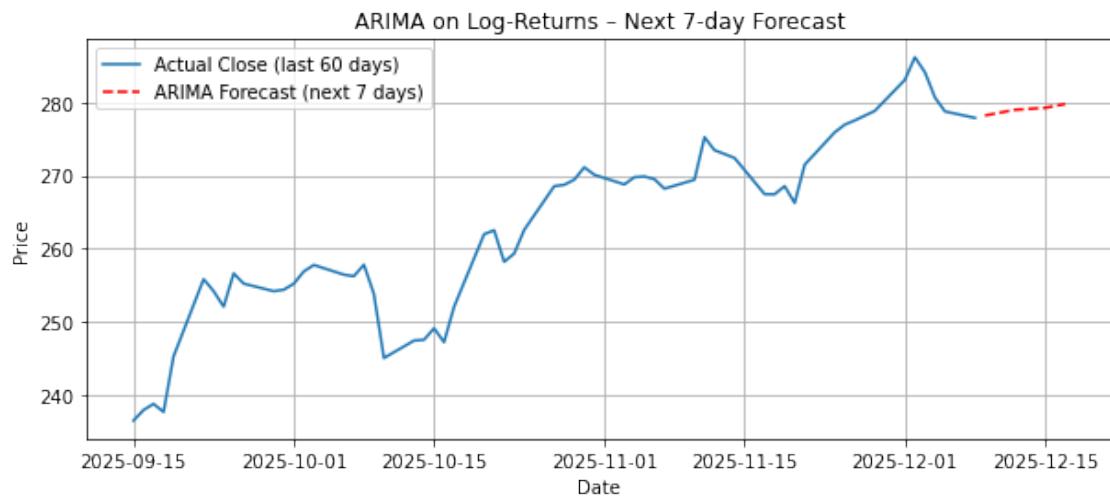
/Users/anujatalekar/opt/anaconda3/lib/python3.9/site-
packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has
been provided, but it has no associated frequency information and so will be
ignored when e.g. forecasting.
    self._init_dates(dates, freq)
/Users/anujatalekar/opt/anaconda3/lib/python3.9/site-
packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has
been provided, but it has no associated frequency information and so will be
ignored when e.g. forecasting.

```

```

    self._init_dates(dates, freq)
/Users/anujatalekar/opt/anaconda3/lib/python3.9/site-
packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: A date index has
been provided, but it has no associated frequency information and so will be
ignored when e.g. forecasting.
    self._init_dates(dates, freq)
/Users/anujatalekar/opt/anaconda3/lib/python3.9/site-
packages/statsmodels/tsa/base/tsa_model.py:834: ValueWarning: No supported index
is available. Prediction results will be given with an integer index beginning
at `start`.
    return get_prediction_index()

```



### 1.1.13 ARIMA 7-Day Forecast Interpretation

ARIMA predicts a stable, slightly upward movement for the next week. The model smooths recent volatility and expects prices to level out around 280.

### 1.1.14 XGBoost Model — Objective, Inputs, and Targets

#### Objective:

To predict the *next-day closing price of AAPL stock* using an advanced machine-learning model capable of learning nonlinear relationships in financial time-series data.

#### Inputs (Features):

The model uses a wide range of engineered features that capture price behavior, momentum, volatility, and historical patterns:

- Raw market variables: **Open, High, Low, Close, Volume**
- Returns-based features: **Return, Log\_Return**
- Trend indicators: **SMA\_10, SMA\_20, SMA\_50**
- Momentum indicators: **EMA\_12, EMA\_26, MACD**
- Overbought/oversold indicator: **RSI**
- Volatility measures: **Vol\_10, Vol\_20**
- Historical dependencies: **Lagged Close values (Lag\_1, Lag\_2, ... Lag\_10)**
- Lagged returns: **Ret\_Lag\_1, Ret\_Lag\_2, Ret\_Lag\_3**

These features allow the model to learn complex interactions that are not captured by simple linear models.

#### Target Variable:

The **actual closing price** of the stock on the same day (used as **Target**). This enables XGBoost to learn a supervised regression mapping between today's features and today's closing price.

### 1.1.15 Motivation for Using XGBoost

Financial markets are nonlinear and influenced by multiple interacting factors such as momentum, volatility, trends, and historical price patterns. Unlike simple time-series models, machine-learning algorithms can capture these complex relationships more effectively.

### 1.1.16 Why XGBoost?

- Handles **nonlinear patterns** in stock price movements that ARIMA cannot model.
- Works well with engineered features such as RSI, MACD, volatility, and lag values.
- Robust to noise and overfitting, which is important for market data.
- Provides **feature importance**, giving insight into which indicators matter most.
- Often achieves **higher predictive performance**, especially in trend and direction forecasting.

XGBoost complements ARIMA by incorporating rich technical indicators and learning from deeper patterns in the data.

```
[247]: raw = pd.read_csv("AAPL_10y_OHLCV.csv", header=None)

#drop first 3 rows
raw = raw.drop(index=[0, 1, 2]).reset_index(drop=True)

#set column names
raw.columns = ["Date", "Close", "High", "Low", "Open", "Volume"]

#convert types
raw["Date"] = pd.to_datetime(raw["Date"], errors="coerce")
num_cols = ["Close", "High", "Low", "Open", "Volume"]
raw[num_cols] = raw[num_cols].apply(pd.to_numeric, errors="coerce")

#remove bad rows, sort, set index
raw = raw.dropna().sort_values("Date").reset_index(drop=True)
df_arima = raw.set_index("Date") # base price dataframe

print(df_arima.head())
print(df_arima.info())
```

```
          Close      High       Low      Open    Volume
Date
2015-12-09  26.068718  26.535439  25.946965  26.524164  185445600
2015-12-10  26.192720  26.366332  26.043912  26.163410  116850800
2015-12-11  25.518581  26.016868  25.444176  25.971775  187544800
2015-12-14  25.360744  25.405838  24.754233  25.293103  257274800
2015-12-15  24.912066  25.432900  24.880500  25.238997  213292400
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2514 entries, 2015-12-09 to 2025-12-08
Data columns (total 5 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   Close    2514 non-null   float64 
 1   High     2514 non-null   float64 
 2   Low      2514 non-null   float64 
 3   Open     2514 non-null   float64 
 4   Volume   2514 non-null   int64  
dtypes: float64(4), int64(1)
memory usage: 117.8 KB
None
```

```
[248]: df_xgb = df_arima.copy()

#returns
df_xgb["Return"] = df_xgb["Close"].pct_change()
```

```

df_xgb["Log_Return"] = np.log(df_xgb["Close"] / df_xgb["Close"].shift(1))

#moving averages
df_xgb["SMA_10"] = df_xgb["Close"].rolling(10).mean()
df_xgb["SMA_20"] = df_xgb["Close"].rolling(20).mean()
df_xgb["SMA_50"] = df_xgb["Close"].rolling(50).mean()

#EMAs & MACD
df_xgb["EMA_12"] = df_xgb["Close"].ewm(span=12, adjust=False).mean()
df_xgb["EMA_26"] = df_xgb["Close"].ewm(span=26, adjust=False).mean()
df_xgb["MACD"] = df_xgb["EMA_12"] - df_xgb["EMA_26"]

#RSI
def compute_rsi(series, period=14):
    delta = series.diff()
    gain = np.where(delta > 0, delta, 0)
    loss = np.where(delta < 0, -delta, 0)
    avg_gain = pd.Series(gain, index=series.index).rolling(period).mean()
    avg_loss = pd.Series(loss, index=series.index).rolling(period).mean()
    rs = avg_gain / avg_loss
    return 100 - (100 / (1 + rs))

df_xgb["RSI"] = compute_rsi(df_xgb["Close"])

#Volatility
df_xgb["Vol_10"] = df_xgb["Close"].rolling(10).std()
df_xgb["Vol_20"] = df_xgb["Close"].rolling(20).std()

#Lags of Close
for k in [1, 2, 3, 5, 7, 10]:
    df_xgb[f'Lag_{k}'] = df_xgb["Close"].shift(k)

#Lags of Return
for k in [1, 2, 3]:
    df_xgb[f'Ret_Lag_{k}'] = df_xgb["Return"].shift(k)

#target = today's Close
df_xgb["Target"] = df_xgb["Close"]

#drop NaNs from rolling/shift
df_xgb = df_xgb.dropna()

print(df_xgb.head())
print("Shape:", df_xgb.shape)

```

	Close	High	Low	Open	Volume	Return	\
Date							
2016-02-22	21.961954	21.966489	21.744330	21.832740	137123200	0.008746	

2016-02-23	21.465504	21.875816	21.433767	21.853147	127770400	-0.022605
2016-02-24	21.785137	21.848611	21.154933	21.304550	145022800	0.014891
2016-02-25	21.934753	21.934753	21.592447	21.773802	110330800	0.006868
2016-02-26	21.968761	22.220388	21.893953	22.034501	115964400	0.001550
Date	Log_Return	SMA_10	SMA_20	SMA_50	...	Lag_1 \
2016-02-22	0.008708	21.669524	21.693281	22.834507	...	21.771540
2016-02-23	-0.022864	21.662270	21.645525	22.742443	...	21.961954
2016-02-24	0.014781	21.687433	21.607550	22.654291	...	21.465504
2016-02-25	0.006844	21.743879	21.651122	22.582615	...	21.785137
2016-02-26	0.001549	21.816647	21.688842	22.514775	...	21.934753
Date	Lag_2	Lag_3	Lag_5	Lag_7	Lag_10	Ret_Lag_1 \
2016-02-22	21.821407	22.243053	21.306814	21.370291	21.313616	-0.002285
2016-02-23	21.771540	21.821407	21.907547	21.241081	21.538044	0.008746
2016-02-24	21.961954	21.771540	22.243053	21.306814	21.533508	-0.022605
2016-02-25	21.465504	21.961954	21.821407	21.907547	21.370291	0.014891
2016-02-26	21.785137	21.465504	21.771540	22.243053	21.241081	0.006868
Date	Ret_Lag_2	Ret_Lag_3	Target			
2016-02-22	-0.018956	0.015315	21.961954			
2016-02-23	-0.002285	-0.018956	21.465504			
2016-02-24	0.008746	-0.002285	21.785137			
2016-02-25	-0.022605	0.008746	21.934753			
2016-02-26	0.014891	-0.022605	21.968761			

[5 rows x 26 columns]

Shape: (2465, 26)

### 1.1.17 Why We Used Technical Indicators, Returns, Lags, and Volatility Features

Stock prices are influenced by several dynamic factors such as momentum, trend strength, volatility, and recent historical patterns.

To capture these complex behaviors, we engineered a diverse set of features commonly used in quantitative finance:

**1. Momentum Indicators (MACD, RSI, EMAs)** These help the model understand whether the stock is gaining upward momentum, slowing down, or entering overbought/oversold regions — patterns that commonly precede reversals or trend continuations.

**2. Trend Indicators (SMAs, EMAs)** Moving averages smooth out noise and help the model identify broader price directions instead of reacting to short-term fluctuations.

**3. Volatility Indicators (Vol\_10, Vol\_20)** High volatility often leads to larger price swings. Including volatility helps the model adjust expectations in turbulent vs. stable periods.

**4. Returns & Log Returns** These stabilize the scale of price movements and make the data more stationary, which helps machine-learning models learn more reliable patterns.

**5. Lag Features (Lag\_1 to Lag\_10)** Stock prices exhibit autocorrelation — today's price is highly dependent on previous days.

Lagged features give the model direct access to recent historical behavior.

### 1.1.18 Overall Benefit

By combining these indicators, the model can learn:

- trends,
- reversals,
- volatility shifts,
- short-term patterns,
- and nonlinear interactions

making XGBoost significantly more expressive and powerful than using raw prices alone.

---

### 1.1.19 Actual vs predicted XGboost

```
[22]: df_xgb_10y = df_xgb.copy()

[23]: from sklearn.preprocessing import StandardScaler
       from xgboost import XGBRegressor
       from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
       import numpy as np
       import matplotlib.pyplot as plt

feature_cols = [
    "Open", "High", "Low", "Volume",
    "Return", "Log_Return",
    "SMA_10", "SMA_20", "SMA_50",
    "EMA_12", "EMA_26",
    "MACD", "RSI",
    "Vol_10", "Vol_20",
    "Lag_1", "Lag_2", "Lag_3", "Lag_5", "Lag_7", "Lag_10",
    "Ret_Lag_1", "Ret_Lag_2", "Ret_Lag_3"
]

X = df_xgb_10y[feature_cols]
y = df_xgb_10y["Target"]

idx = df_xgb_10y.index

#test period = 2023 + 2024
test_mask = (idx >= "2023-01-01") & (idx <= "2024-12-31")
```

```

#train on everything else (2016-2022 and 2025+ inside the 10y slice)
train_mask = ~test_mask

X_train, X_test = X[train_mask], X[test_mask]
y_train, y_test = y[train_mask], y[test_mask]

print("Train rows:", len(X_train), "Test rows:", len(X_test))

#scale
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

model_xgb = XGBRegressor(
    n_estimators=1200,
    learning_rate=0.02,
    max_depth=7,
    subsample=0.9,
    colsample_bytree=0.9,
    reg_lambda=0.5,
    objective="reg:squarederror",
    random_state=42,
)
model_xgb.fit(X_train_scaled, y_train)

#predictions on full 10y window
X_all_scaled = scaler.transform(X)
pred_all = model_xgb.predict(X_all_scaled)

# metrics on 2023-2024 only
pred_test = model_xgb.predict(X_test_scaled)
rmse = np.sqrt(mean_squared_error(y_test, pred_test))
mae = mean_absolute_error(y_test, pred_test)
r2 = r2_score(y_test, pred_test)
mape = np.mean(np.abs((y_test - pred_test) / y_test)) * 100
accuracy = 100 - mape

print("XGBoost (10y, train = all except 2023-2024, test = 2023-2024) ↴")
print("RMSE:", rmse)
print("MAE :", mae)
print("R²  :", r2)
print("MAPE:", mape)
print("Accuracy (%):", accuracy)

#plot full 10y

```

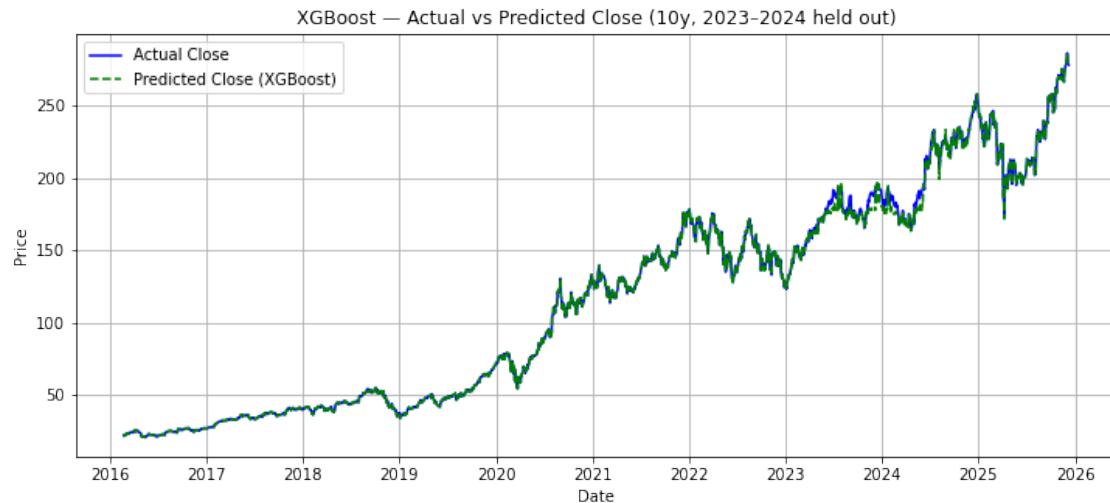
```

dates_10y = df_xgb_10y.index
actual_10y = df_xgb_10y["Close"].values

plt.figure(figsize=(12,5))
plt.plot(dates_10y, actual_10y, label="Actual Close", color="blue")
plt.plot(dates_10y, pred_all, label="Predicted Close (XGBoost)",
         color="green", linestyle="--")
plt.title("XGBoost - Actual vs Predicted Close (10y, 2023-2024 held out)")
plt.xlabel("Date")
plt.ylabel("Price")
plt.grid(True)
plt.legend()
plt.show()

```

Train rows: 1963 Test rows: 502  
 XGBoost (10y, train = all except 2023-2024, test = 2023-2024) performance:  
 RMSE: 4.581475167366499  
 MAE : 2.9602649582334717  
 $R^2$  : 0.9733912102951328  
 MAPE: 1.5633912127156622  
 Accuracy (%): 98.43660878728434



### 1.1.20 Trying to find the feature importance

```
[24]: import numpy as np
import pandas as pd

#get raw importances from XGBoost
importances = model_xgb.feature_importances_
```

```

#put into a DataFrame for easier handling
fi_df = pd.DataFrame({
    "feature": feature_cols,
    "importance": importances
})

#sort descending
fi_df = fi_df.sort_values("importance", ascending=False).reset_index(drop=True)

#compute percentage importance
fi_df["importance_pct"] = fi_df["importance"] / fi_df["importance"].sum() * 100

print(fi_df)
print("\nTop 10 features:")
print(fi_df.head(10))

```

	feature	importance	importance_pct
0	High	0.259175	25.917534
1	Open	0.243086	24.308643
2	SMA_10	0.207047	20.704693
3	Low	0.134382	13.438227
4	Lag_2	0.093839	9.383862
5	Lag_1	0.028380	2.837968
6	SMA_50	0.013501	1.350067
7	Lag_7	0.010296	1.029583
8	SMA_20	0.005620	0.562043
9	EMA_12	0.002481	0.248091
10	Lag_3	0.000926	0.092625
11	Lag_5	0.000589	0.058916
12	Lag_10	0.000443	0.044327
13	EMA_26	0.000129	0.012881
14	Log_Return	0.000036	0.003577
15	Return	0.000031	0.003075
16	RSI	0.000007	0.000658
17	Vol_10	0.000006	0.000575
18	Ret_Lag_1	0.000006	0.000565
19	Vol_20	0.000005	0.000531
20	Volume	0.000004	0.000421
21	MACD	0.000004	0.000405
22	Ret_Lag_2	0.000004	0.000383
23	Ret_Lag_3	0.000004	0.000357

Top 10 features:

	feature	importance	importance_pct
0	High	0.259175	25.917534
1	Open	0.243086	24.308643
2	SMA_10	0.207047	20.704693
3	Low	0.134382	13.438227

```

4   Lag_2      0.093839      9.383862
5   Lag_1      0.028380      2.837968
6   SMA_50     0.013501      1.350067
7   Lag_7      0.010296      1.029583
8   SMA_20     0.005620      0.562043
9   EMA_12     0.002481      0.248091

```

```

[25]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

#get feature importances from the trained XGBoost model
importances = model_xgb.feature_importances_

fi_df = pd.DataFrame({
    "feature": feature_cols,
    "importance": importances
})

#sort and turn into percentages
fi_df = fi_df.sort_values("importance", ascending=False).reset_index(drop=True)
fi_df["importance_pct"] = fi_df["importance"] / fi_df["importance"].sum() * 100

print("Top 10 features by importance:")
print(fi_df.head(10))

#plot top N features
N = 15
topN = fi_df.head(N)

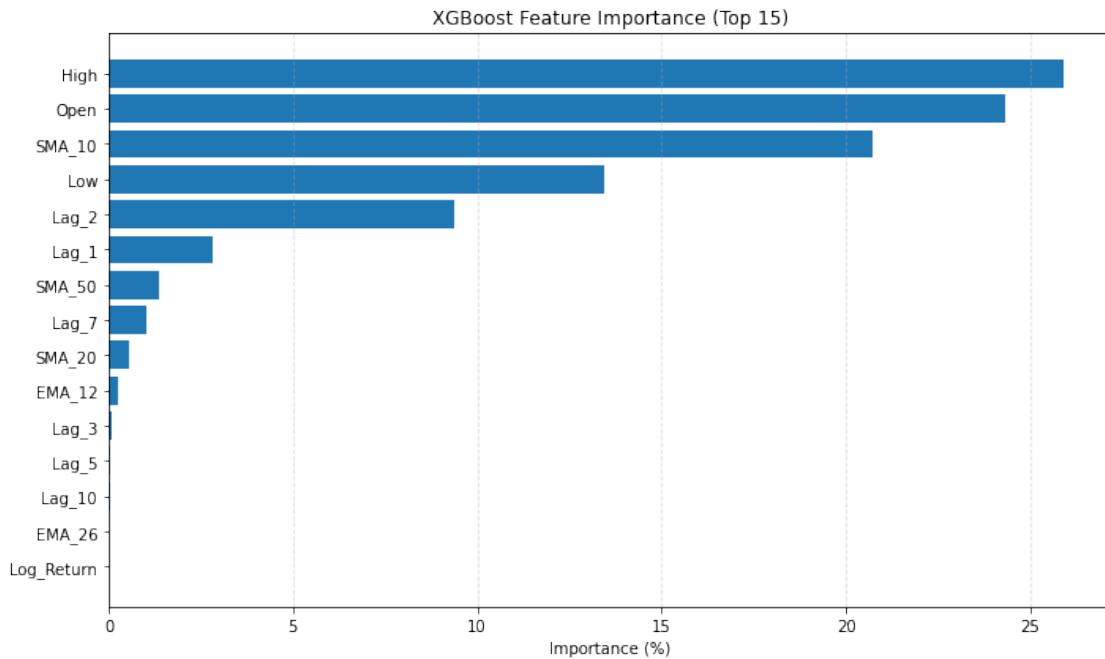
plt.figure(figsize=(10, 6))
plt.barh(topN["feature"], topN["importance_pct"])
plt.gca().invert_yaxis() # most important at the top
plt.xlabel("Importance (%)")
plt.title("XGBoost Feature Importance (Top 15)")
plt.grid(axis="x", linestyle="--", alpha=0.4)
plt.tight_layout()
plt.show()

```

Top 10 features by importance:

	feature	importance	importance_pct
0	High	0.259175	25.917534
1	Open	0.243086	24.308643
2	SMA_10	0.207047	20.704693
3	Low	0.134382	13.438227
4	Lag_2	0.093839	9.383862
5	Lag_1	0.028380	2.837968
6	SMA_50	0.013501	1.350067

7	Lag_7	0.010296	1.029583
8	SMA_20	0.005620	0.562043
9	EMA_12	0.002481	0.248091



### 1.1.21 Importance of Feature Engineering in Stock Price Prediction

Financial markets are **nonlinear and highly noisy**, meaning raw price data alone does not provide enough structure for a machine-learning model to learn meaningful patterns.

Feature engineering allows us to extract financial intuition and embed it directly into the model through:

- **Momentum** (RSI, MACD, EMAs)
- **Trend strength** (SMAs)
- **Volatility regimes** (rolling standard deviation)
- **Historical influence** (lag features)
- **Normalized changes** (returns)

This enriched feature set converts raw chaotic price movements into **interpretable signals**, enabling XGBoost to outperform traditional statistical models like ARIMA, especially in capturing nonlinear market behavior.

### 1.1.22 Next 7 day projection for XGBoost

```
[27]: import pandas as pd
import numpy as np
from pandas.tseries.offsets import BDay

# ----- helper: recompute all features from price df -----
def make_features_from_price(df_price):
    df = df_price.copy()

    # basic returns
    df["Return"] = df["Close"].pct_change()
    df["Log_Return"] = np.log(df["Close"] / df["Close"].shift(1))

    # moving averages
    df["SMA_10"] = df["Close"].rolling(10).mean()
    df["SMA_20"] = df["Close"].rolling(20).mean()
    df["SMA_50"] = df["Close"].rolling(50).mean()

    # EMAs & MACD
    df["EMA_12"] = df["Close"].ewm(span=12, adjust=False).mean()
    df["EMA_26"] = df["Close"].ewm(span=26, adjust=False).mean()
    df["MACD"] = df["EMA_12"] - df["EMA_26"]

    # RSI
    def compute_rsi(series, period=14):
        delta = series.diff()
        gain = np.where(delta > 0, delta, 0)
        loss = np.where(delta < 0, -delta, 0)
        avg_gain = pd.Series(gain, index=series.index).rolling(period).mean()
        avg_loss = pd.Series(loss, index=series.index).rolling(period).mean()
        rs = avg_gain / (avg_loss + 1e-9) # avoid div/0
        return 100 - (100 / (1 + rs))

    df["RSI"] = compute_rsi(df["Close"])

    # volatility
    df["Vol_10"] = df["Close"].rolling(10).std()
    df["Vol_20"] = df["Close"].rolling(20).std()

    # lags of Close
    for k in [1, 2, 3, 5, 7, 10]:
        df[f"Lag_{k}"] = df["Close"].shift(k)

    # lags of Return
    for k in [1, 2, 3]:
        df[f"Ret_Lag_{k}"] = df["Return"].shift(k)
```

```

df = df.dropna()
return df

# ----- start from last 10 years of price data -----
price_hist = df_arima.last("10Y").copy()

future_dates = []
future_preds_xgb = []

steps_ahead = 7

for step in range(steps_ahead):
    # 1) Recompute features on current history
    feat_df = make_features_from_price(price_hist)
    X_last = feat_df[feature_cols].iloc[-1:] # last row

    # 2) Scale + predict next-day Close
    X_last_scaled = scaler.transform(X_last)
    next_close = model_xgb.predict(X_last_scaled)[0]

    # 3) Decide next date (next business day)
    last_date = price_hist.index[-1]
    next_date = last_date + BDay(1)

    future_dates.append(next_date)
    future_preds_xgb.append(next_close)

    # 4) Append synthetic OHLCV row for the predicted day
    #     (simple assumption: Open=High=Low=Close, Volume same as last day)
    new_row = pd.DataFrame({
        "Open": [next_close],
        "High": [next_close],
        "Low": [next_close],
        "Close": [next_close],
        "Volume": [price_hist["Volume"].iloc[-1]]
    }, index=[next_date])

    price_hist = pd.concat([price_hist, new_row])

# Collect results
future_xgb_df = pd.DataFrame(
    {"Predicted_Close_XGB": future_preds_xgb},
    index=future_dates
)

print("Next 7-day XGBoost forecast:")

```

```

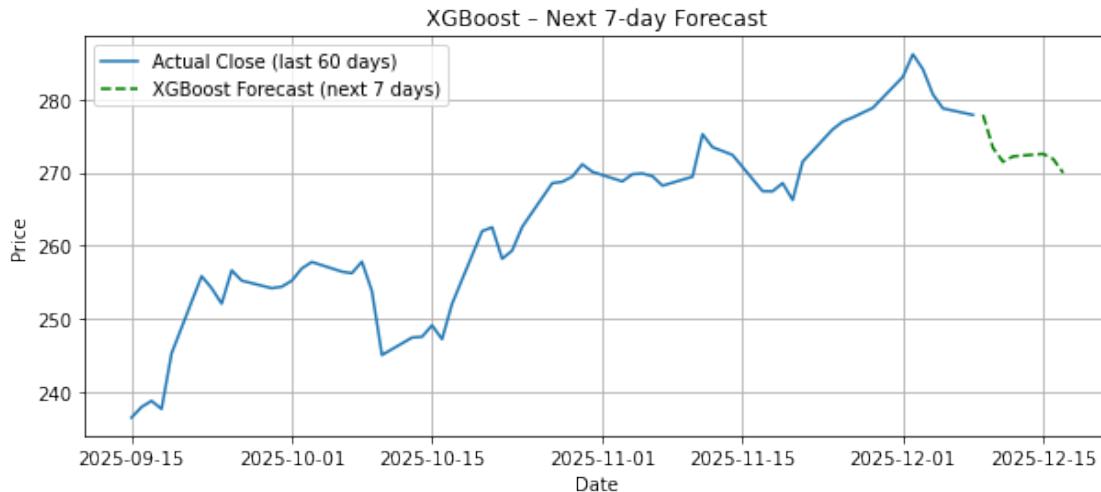
print(future_xgb_df)

# Optional: plot last 60 days + 7-day XGBoost forecast
plt.figure(figsize=(10,4))
plt.plot(df_arima["Close"].tail(60), label="Actual Close (last 60 days)")
plt.plot(future_xgb_df["Predicted_Close_XGB"], label="XGBoost Forecast (next 7 days)",
         linestyle="--", color="green")
plt.xlabel("Date")
plt.ylabel("Price")
plt.title("XGBoost - Next 7-day Forecast")
plt.grid(True)
plt.legend()
plt.show()

```

Next 7-day XGBoost forecast:

	Predicted_Close_XGB
2025-12-09	277.888733
2025-12-10	273.365051
2025-12-11	271.478699
2025-12-12	272.214844
2025-12-15	272.562103
2025-12-16	271.867065
2025-12-17	269.972748

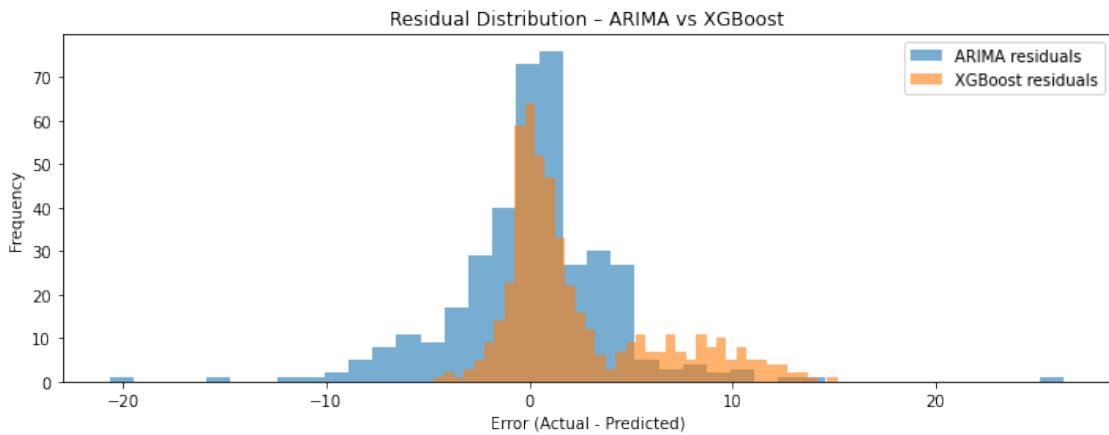


### 1.1.23 XGBoost 7-Day Forecast Interpretation

XGBoost predicts a mild short-term decline, dropping from ~278 to ~270 over the next week. It reacts to the recent dip in prices and expects continued soft downward movement.

### 1.1.24 Residual Disrtribution

```
[36]: plt.figure(figsize=(10,4))
plt.hist(err_arima, bins=40, alpha=0.6, label="ARIMA residuals")
plt.hist(err_xgb, bins=40, alpha=0.6, label="XGBoost residuals")
plt.title("Residual Distribution - ARIMA vs XGBoost")
plt.xlabel("Error (Actual - Predicted)")
plt.ylabel("Frequency")
plt.legend()
plt.tight_layout()
plt.show()
```



### Residual distribution – ARIMA vs XGBoost

- X-axis: forecasting error = Actual – Predicted
  - 0 perfect prediction
  - Positive error model **under-predicts** (actual > predicted)
  - Negative error model **over-predicts** (actual < predicted)
- Y-axis: how often each error size occurs.

From the histograms: - **ARIMA residuals (blue)** are wider and spread on both sides → larger and more variable errors (both big over- and under-predictions).

- **XGBoost residuals (orange)** are more concentrated but slightly shifted to the right → it tends to **under-predict** prices a bit, but with generally smaller error variance than ARIMA.

## 1.2 SVR Model

### 1.2.1 Inputs & Targets

**INPUTS:** - Historical **OHLCV** data (Open, High, Low, Close, Volume)  
- Engineered **technical indicators** (moving averages, volatility, momentum)

- **Temporal features** (day of week, month)

**TARGETS:** - **Weekly forward returns** (% change in closing price)

---

### 1.2.2 Motivation

Stock price prediction is crucial for **portfolio management** and **risk assessment**.

Support Vector Regression (**SVR**) is particularly effective for this task because:

1. **Captures non-linear relationships** in financial data
2. **Robust to outliers** via epsilon-insensitive loss
3. **Kernel trick** allows detection of complex patterns
4. **Regularization** helps prevent overfitting in noisy market environments

```
[240]: """
SVR DAILY STOCK PRICE PREDICTION MODEL

"""

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.svm import SVR
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from scipy import stats
import warnings
import os
import pickle

warnings.filterwarnings('ignore')
os.makedirs('outputs', exist_ok=True)

print("SVR DAILY STOCK PRICE PREDICTION MODEL (SAFE VERSION)")

# 1. DATA ACQUISITION & PREPARATION

raw_svr_daily = pd.read_csv("AAPL_since_IPO_OHLCV.csv", header=None)
raw_svr_daily = raw_svr_daily.drop(index=[0, 1, 2]).reset_index(drop=True)
raw_svr_daily.columns = ["Date", "Close", "High", "Low", "Open", "Volume"]

raw_svr_daily["Date"] = pd.to_datetime(raw_svr_daily["Date"], errors="coerce")
num_cols_svr_daily = ["Close", "High", "Low", "Open", "Volume"]
```

```

raw_svr_daily[num_cols_svr_daily] = raw_svr_daily[num_cols_svr_daily].apply(pd.
    ↪to_numeric, errors="coerce")
raw_svr_daily = raw_svr_daily.dropna().sort_values("Date").
    ↪reset_index(drop=True)

recent_cutoff_svr_daily = raw_svr_daily["Date"].max() - pd.
    ↪Timedelta(days=365*10)
df_svr_daily = raw_svr_daily[raw_svr_daily["Date"] >= recent_cutoff_svr_daily].
    ↪copy()
df_svr_daily = df_svr_daily.set_index("Date")

# 2. FEATURE ENGINEERING

df_svr_daily["Daily_Return"] = df_svr_daily["Close"].pct_change()
df_svr_daily["SMA_5"] = df_svr_daily["Close"].rolling(5).mean()
df_svr_daily["SMA_20"] = df_svr_daily["Close"].rolling(20).mean()
df_svr_daily["SMA_50"] = df_svr_daily["Close"].rolling(50).mean()

df_svr_daily["EMA_12"] = df_svr_daily["Close"].ewm(span=12, adjust=False).mean()
df_svr_daily["EMA_26"] = df_svr_daily["Close"].ewm(span=26, adjust=False).mean()
df_svr_daily["MACD"] = df_svr_daily["EMA_12"] - df_svr_daily["EMA_26"]

df_svr_daily["Volatility"] = df_svr_daily["Daily_Return"].rolling(20).std()
df_svr_daily["Momentum_5"] = df_svr_daily["Close"] - df_svr_daily["Close"].
    ↪shift(5)
df_svr_daily["Momentum_20"] = df_svr_daily["Close"] - df_svr_daily["Close"].
    ↪shift(20)

delta = df_svr_daily["Close"].diff()
gain = (delta.where(delta > 0, 0)).rolling(14).mean()
loss = (-delta.where(delta < 0, 0)).rolling(14).mean()
rs = gain / loss
df_svr_daily["RSI"] = 100 - (100 / (1 + rs))

df_svr_daily["BB_middle"] = df_svr_daily["Close"].rolling(20).mean()
df_svr_daily["BB_std"] = df_svr_daily["Close"].rolling(20).std()
df_svr_daily["BB_upper"] = df_svr_daily["BB_middle"] + 2 * ↪
    ↪df_svr_daily["BB_std"]
df_svr_daily["BB_lower"] = df_svr_daily["BB_middle"] - 2 * ↪
    ↪df_svr_daily["BB_std"]
df_svr_daily["BB_width"] = df_svr_daily["BB_upper"] - df_svr_daily["BB_lower"]

df_svr_daily["Volume_MA_20"] = df_svr_daily["Volume"].rolling(20).mean()
df_svr_daily["Volume_Ratio"] = df_svr_daily["Volume"] / ↪
    ↪df_svr_daily["Volume_MA_20"]

```

```

df_svr_daily["High_Low_Range"] = df_svr_daily["High"] - df_svr_daily["Low"]
df_svr_daily["Close_Open_Range"] = df_svr_daily["Close"] - df_svr_daily["Open"]

for lag in [1, 2, 3, 5, 10]:
    df_svr_daily[f"Return_Lag_{lag}"] = df_svr_daily["Daily_Return"].shift(lag)

df_svr_daily["DayOfWeek"] = df_svr_daily.index.dayofweek
df_svr_daily["Month"] = df_svr_daily.index.month

df_svr_daily["Target"] = df_svr_daily["Close"].shift(-1)
df_svr_daily = df_svr_daily.dropna()

# 3. TRAIN-TEST SPLIT + SCALING

feature_columns_svr_daily = [col for col in df_svr_daily.columns if col not in
                             ["Close", "Open", "High", "Low", "Daily_Return", ↴
                             "Target"]]

X_svr_daily = df_svr_daily[feature_columns_svr_daily]
y_svr_daily = df_svr_daily["Target"]

split_idx = int(len(X_svr_daily) * 0.8)
X_train_svr_daily = X_svr_daily.iloc[:split_idx]
X_test_svr_daily = X_svr_daily.iloc[split_idx:]
y_train_svr_daily = y_svr_daily.iloc[:split_idx]
y_test_svr_daily = y_svr_daily.iloc[split_idx:]

scaler_svr_daily = StandardScaler()
X_train_svr_daily_scaled = scaler_svr_daily.fit_transform(X_train_svr_daily)
X_test_svr_daily_scaled = scaler_svr_daily.transform(X_test_svr_daily)

# 4. SVR MODELS

results_svr_daily = []
tscv = TimeSeriesSplit(n_splits=5)

def evaluate_model(name, model):
    model.fit(X_train_svr_daily_scaled, y_train_svr_daily)
    pred_train = model.predict(X_train_svr_daily_scaled)
    pred_test = model.predict(X_test_svr_daily_scaled)

    cv_rmse = []
    for tr, val in tscv.split(X_train_svr_daily_scaled):
        m = model.__class__(**model.get_params())
        m.fit(X_train_svr_daily_scaled[tr], y_train_svr_daily.iloc[tr])
        val_pred = m.predict(X_train_svr_daily_scaled[val])

```

```

        cv_rmse.append(np.sqrt(mean_squared_error(y_train_svr_daily.iloc[val], val_pred)))

    results_svr_daily[name] = {
        "model": model,
        "train_rmse": np.sqrt(mean_squared_error(y_train_svr_daily, pred_train)),
        "test_rmse": np.sqrt(mean_squared_error(y_test_svr_daily, pred_test)),
        "test_mae": mean_absolute_error(y_test_svr_daily, pred_test),
        "test_r2": r2_score(y_test_svr_daily, pred_test),
        "cv_rmse": np.mean(cv_rmse),
        "pred_test": pred_test,
    }

# Train models
evaluate_model("Linear", SVR(kernel="linear", C=1, epsilon=0.1))
evaluate_model("RBF", SVR(kernel="rbf", C=1, epsilon=0.1, gamma="scale"))
evaluate_model("Polynomial", SVR(kernel="poly", C=1, epsilon=0.1, degree=3))

# 5. SELECT BEST MODEL

best_svr_daily_model_name = min(results_svr_daily, key=lambda k: results_svr_daily[k]["test_rmse"])
best_svr_daily_model = results_svr_daily[best_svr_daily_model_name]["model"]

print(f"\nBEST DAILY SVR MODEL: {best_svr_daily_model_name}")

# 6. 7-DAY FORECAST

future_predictions_svr_daily = []
future_dates_svr_daily = []

last_row = X_svr_daily.iloc[-1:]
last_date = X_svr_daily.index[-1]

for d in range(1, 8):
    scaled = scaler_svr_daily.transform(last_row)
    next_price = best_svr_daily_model.predict(scaled)[0]
    future_predictions_svr_daily.append(next_price)

    next_date = last_date + pd.Timedelta(days=d)
    while next_date.dayofweek >= 5:
        next_date += pd.Timedelta(days=1)
    future_dates_svr_daily.append(next_date)

print("\n7-Day SVR DAILY FORECAST COMPLETE.")

```

## SVR DAILY STOCK PRICE PREDICTION MODEL (SAFE VERSION)

BEST DAILY SVR MODEL: Linear

7-Day SVR DAILY FORECAST COMPLETE.

```
[241]: # SVR DAILY - EVALUATION & VISUALIZATIONS (SAFE)
```

### # 1. METRICS COMPARISON TABLE

```
comparison_df_svr_daily = pd.DataFrame({
    "Model": list(results_svr_daily.keys()),
    "Train_RMSE": [results_svr_daily[m] ["train_rmse"] for m in ↵
    ↵results_svr_daily],
    "Test_RMSE": [results_svr_daily[m] ["test_rmse"] for m in ↵
    ↵results_svr_daily],
    "Test_MAE": [results_svr_daily[m] ["test_mae"] for m in ↵
    ↵results_svr_daily],
    "Test_R2": [results_svr_daily[m] ["test_r2"] for m in ↵
    ↵results_svr_daily],
    "CV_RMSE": [results_svr_daily[m] ["cv_rmse"] for m in ↵
    ↵results_svr_daily],
})
print("\nSVR DAILY - Performance Comparison (Train/Test/CV):")
print(comparison_df_svr_daily.to_string(index=False))

print(f"\nBest model (already selected): {best_svr_daily_model_name}")
best_pred_test_svr_daily = ↵
    ↵results_svr_daily[best_svr_daily_model_name] ["pred_test"]
```

### # 2. ERROR METRICS FOR BEST MODEL

```
rmse_best = results_svr_daily[best_svr_daily_model_name] ["test_rmse"]
mae_best = results_svr_daily[best_svr_daily_model_name] ["test_mae"]
r2_best = results_svr_daily[best_svr_daily_model_name] ["test_r2"]

mape_best = np.mean(
    np.abs((y_test_svr_daily.values - best_pred_test_svr_daily) / ↵
    ↵y_test_svr_daily.values)
) * 100
acc_best = 100 - mape_best

print("\nDETAILED METRICS - BEST SVR DAILY MODEL")
print("-"*80)
print(f"Model: {best_svr_daily_model_name} kernel")
print(f"Test RMSE : {rmse_best:.4f}")
```

```

print(f"Test MAE : {mae_best:.4f}")
print(f"Test R2 : {r2_best:.4f}")
print(f"MAPE : {mape_best:.2f}%")
print(f"Accuracy : {acc_best:.2f}%")

# 3. ACTUAL vs PREDICTED (TEST SET) PLOT

dates_test_svr_daily = y_test_svr_daily.index
actual_test_svr_daily = y_test_svr_daily.values

plt.figure(figsize=(12,5))
plt.plot(dates_test_svr_daily, actual_test_svr_daily,
          label="Actual Close (Test)", color="blue", linewidth=1.5)
plt.plot(dates_test_svr_daily, best_pred_test_svr_daily,
          label=f"Predicted Close - SVR ({best_svr_daily_model_name})",
          color="red", linestyle="--", linewidth=1.5)
plt.title(f"SVR DAILY ({best_svr_daily_model_name}) - Actual vs Predicted (Test Set)")
plt.xlabel("Date")
plt.ylabel("Price ($)")
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

# 4. RESIDUAL ANALYSIS (TIME SERIES + HISTOGRAM)

residuals_svr_daily = actual_test_svr_daily - best_pred_test_svr_daily

print("\nRESIDUAL STATS - BEST SVR DAILY MODEL")
print("-"*80)
print(f"Mean residual : {np.mean(residuals_svr_daily):.4f}")
print(f"Std residual : {np.std(residuals_svr_daily, ddof=1):.4f}")
print(f"Skewness : {stats.skew(residuals_svr_daily):.4f}")
print(f"Kurtosis (excess) : {stats.kurtosis(residuals_svr_daily):.4f}")

# Normality test
stat_norm, p_norm = stats.normaltest(residuals_svr_daily)
print(f"\nNormality test (D'Agostino-Pearson): stat={stat_norm:.3f}, p={p_norm:.4f}")
if p_norm > 0.05:
    print(" → Residuals are roughly consistent with normality (p > 0.05)")
else:
    print(" → Residuals deviate from normality (p ≤ 0.05)")

# Residual over time
plt.figure(figsize=(12,4))

```

```

plt.plot(dates_test_svr_daily, residuals_svr_daily, label="Residuals (Actual - Predicted)")
plt.axhline(0, color="black", linewidth=1)
plt.title(f"SVR DAILY ({best_svr_daily_model_name}) - Residuals over Time (Test)")
plt.xlabel("Date")
plt.ylabel("Error ($)")
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

# Residual histogram
plt.figure(figsize=(8,4))
plt.hist(residuals_svr_daily, bins=40, alpha=0.7, edgecolor="black")
plt.title(f"SVR DAILY ({best_svr_daily_model_name}) - Residual Distribution (Test)")
plt.xlabel("Residual (Actual - Predicted)")
plt.ylabel("Frequency")
plt.grid(True, axis="y", alpha=0.3)
plt.tight_layout()
plt.show()

# 5. DIRECTIONAL ACCURACY

# previous day's close for the test period
prev_close_svr_daily = df_svr_daily.loc[y_test_svr_daily.index, "Close"].
    shift(1)

direction_actual_svr_daily = np.sign(actual_test_svr_daily - prev_close_svr_daily.values)
direction_pred_svr_daily = np.sign(best_pred_test_svr_daily - prev_close_svr_daily.values)

mask_dir = ~np.isnan(direction_actual_svr_daily)
dir_acc_svr_daily = (direction_actual_svr_daily[mask_dir] ==
                     direction_pred_svr_daily[mask_dir]).mean() * 100

print("\nDIRECTIONAL ACCURACY - BEST SVR DAILY MODEL")
print("-"*80)
print(f"Directional accuracy: {dir_acc_svr_daily:.2f}%")
print("(% of times model correctly predicts UP/DOWN move from previous close)")

# 6. CONFIDENCE-BAND STYLE SUMMARY (SIMPLE)

residual_std_svr_daily = np.std(residuals_svr_daily, ddof=1)

```

```

band_width = 1.96 * residual_std_svr_daily
print("\nAPPROX 95% PREDICTION BAND (EMPIRICAL)")
print("-"*80)
print(f"Approx ± band width: ±{band_width:.4f} USD around prediction")

# 7. BAR CHARTS: RMSE & R² ACROSS KERNELS

models_names = list(results_svr_daily.keys())
rmse_vals = [results_svr_daily[m]["test_rmse"] for m in models_names]
r2_vals = [results_svr_daily[m]["test_r2"] for m in models_names]

colors_bar = ["#2ecc71" if m == best_svr_daily_model_name else "#95a5a6"
              for m in models_names]

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(13,5))

# RMSE bar chart
ax1.bar(models_names, rmse_vals, color=colors_bar, edgecolor="black")
ax1.set_title("SVR DAILY - Test RMSE by Kernel")
ax1.set_ylabel("RMSE ($)")
ax1.grid(True, axis="y", alpha=0.3)
for i, v in enumerate(rmse_vals):
    ax1.text(i, v, f"{v:.2f}", ha="center", va="bottom", fontweight="bold")

# R² bar chart
ax2.bar(models_names, r2_vals, color=colors_bar, edgecolor="black")
ax2.set_title("SVR DAILY - Test R² by Kernel")
ax2.set_ylabel("R² Score")
ax2.axhline(0, color="red", linestyle="--", alpha=0.5)
ax2.grid(True, axis="y", alpha=0.3)
for i, v in enumerate(r2_vals):
    ax2.text(i, v, f"{v:.3f}", ha="center", va="bottom", fontweight="bold")

plt.tight_layout()
plt.show()

print("\nEVALUATION & VISUALIZATION FOR SVR DAILY MODEL COMPLETE.")

```

SVR DAILY - Performance Comparison (Train/Test/CV):

Model	Train_RMSE	Test_RMSE	Test_MAE	Test_R2	CV_RMSE
Linear	2.006397	4.139530	2.892753	0.977881	2.110628
RBF	11.979452	80.769448	70.989157	-7.420940	34.021039
Polynomial	20.378585	298.491881	231.102655	-114.008884	28.992204

Best model (already selected): Linear

## DETAILED METRICS - BEST SVR DAILY MODEL

---

Model: Linear kernel

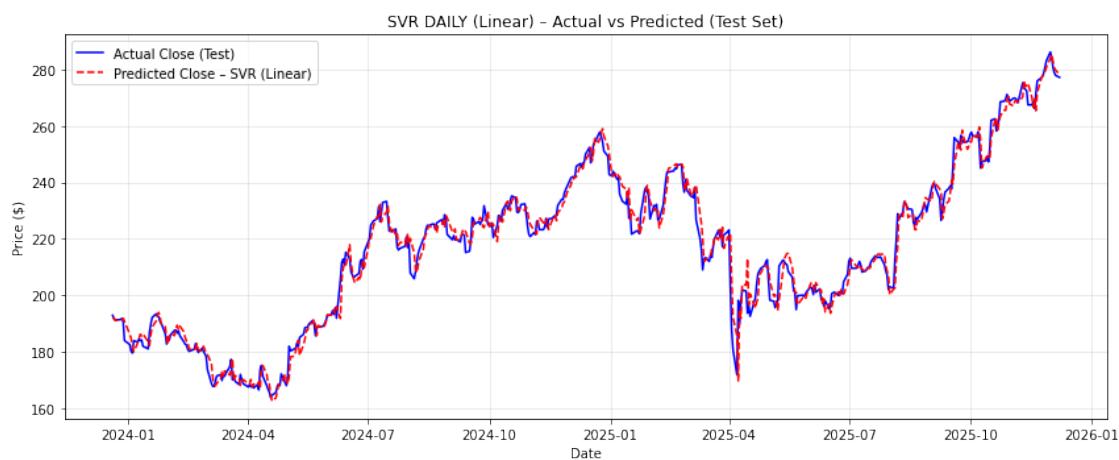
Test RMSE : 4.1395

Test MAE : 2.8928

Test R<sup>2</sup> : 0.9779

MAPE : 1.36%

Accuracy : 98.64%



## RESIDUAL STATS - BEST SVR DAILY MODEL

---

Mean residual : -0.1409

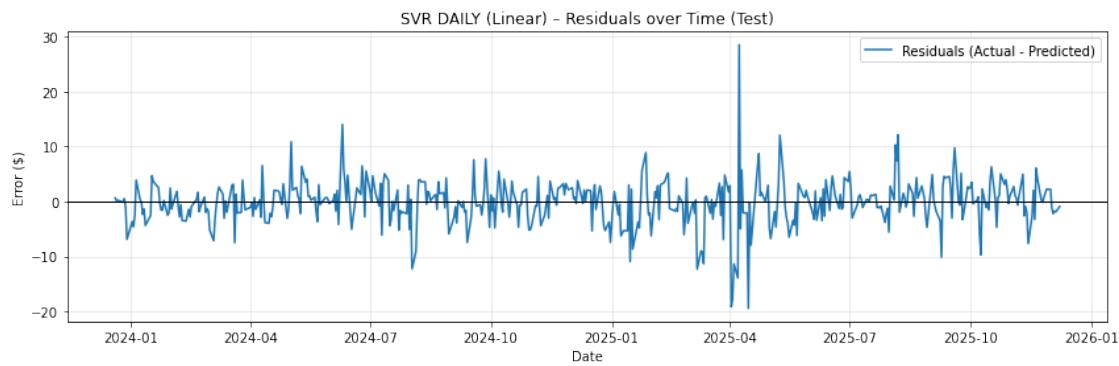
Std residual : 4.1413

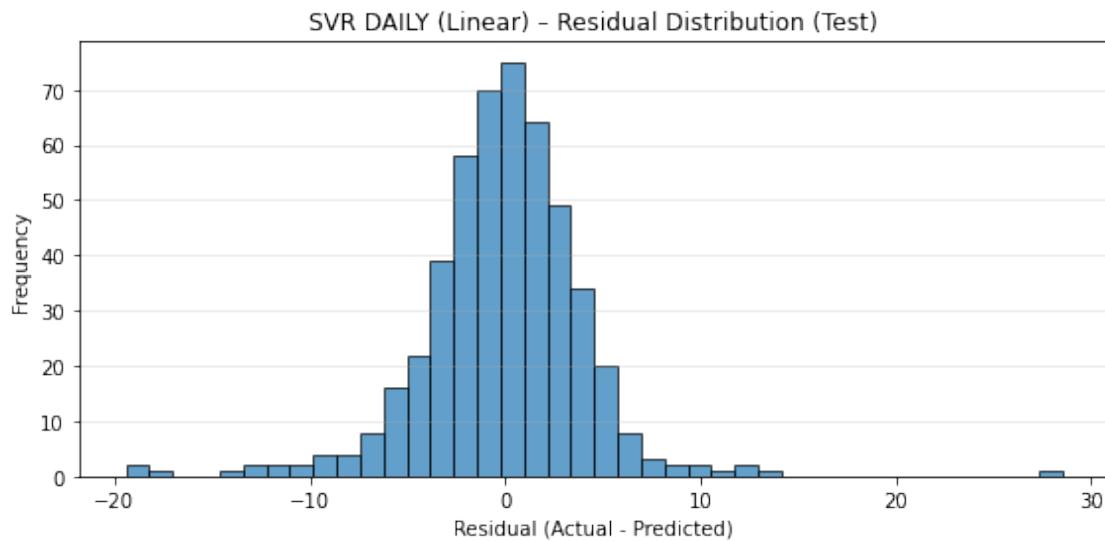
Skewness : 0.0344

Kurtosis (excess) : 6.9880

Normality test (D'Agostino-Pearson): stat=78.913, p=0.0000

→ Residuals deviate from normality (p < 0.05)





#### DIRECTIONAL ACCURACY - BEST SVR DAILY MODEL

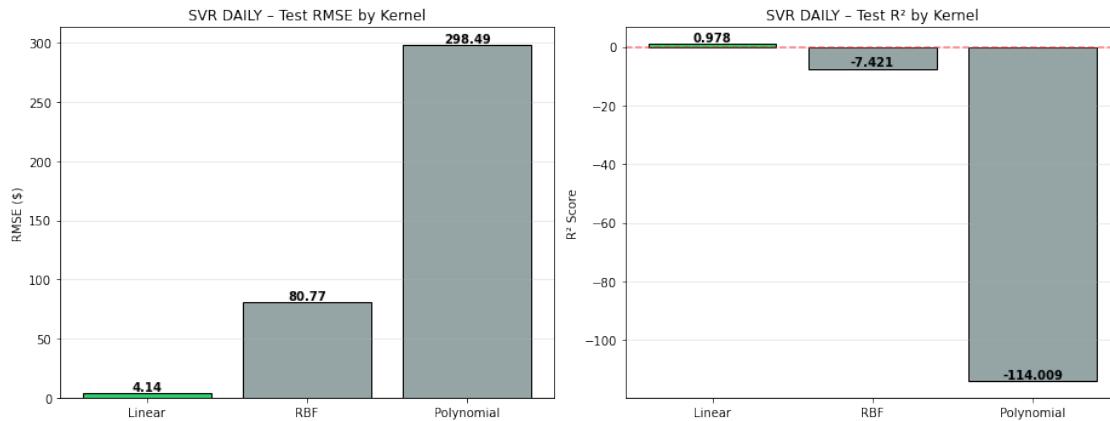
---

Directional accuracy: 70.33%  
 (% of times model correctly predicts UP/DOWN move from previous close)

#### APPROX 95% PREDICTION BAND (EMPIRICAL)

---

Approx  $\pm$  band width:  $\pm 8.1170$  USD around prediction



EVALUATION & VISUALIZATION FOR SVR DAILY MODEL COMPLETE.

### 1.2.3 SVR-A Next seven days projection

```
[242]: # SVR-A (Daily SVR) - 7-DAY PRICE FORECAST VISUALIZATION

# 1. GENERATE NEXT 7 TRADING DAYS USING ROLLING PREDICTIONS


future_predictions_svrA = []
future_dates_svrA = []

# Start from the last available row of features
current_features = X_svr_daily.iloc[-1:].copy()
last_date = X_svr_daily.index[-1]

for step in range(1, 8):
    # scale features
    scaled_features = scaler_svr_daily.transform(current_features)

    # predict next day's closing price
    next_price = best_svr_daily_model.predict(scaled_features)[0]

    # save prediction
    future_predictions_svrA.append(next_price)

    # move date forward (skip weekends)
    next_date = last_date + pd.Timedelta(days=1)
    while next_date.dayofweek >= 5:
        next_date += pd.Timedelta(days=1)
    future_dates_svrA.append(next_date)

    # create synthetic next-day row for rolling forecasting
    temp = current_features.copy()
    temp.iloc[0] = temp.iloc[0] # keep structure
    temp.index = [next_date]

    # Insert predicted Close price into new OHLCV row
    new_row = {
        "Close": next_price,
        "High": next_price * 1.002,
        "Low": next_price * 0.998,
        "Open": next_price,
        "Volume": df_svr_daily["Volume"].iloc[-1],
    }
    df_temp = df_svr_daily.append(pd.DataFrame([new_row], index=[next_date]))

    # Recompute all features dynamically for next step
    temp_full = df_temp.copy()
```

```

temp_full["Daily_Return"] = temp_full["Close"].pct_change()
temp_full["SMA_5"] = temp_full["Close"].rolling(5).mean()
temp_full["SMA_20"] = temp_full["Close"].rolling(20).mean()
temp_full["SMA_50"] = temp_full["Close"].rolling(50).mean()

temp_full["EMA_12"] = temp_full["Close"].ewm(span=12, adjust=False).mean()
temp_full["EMA_26"] = temp_full["Close"].ewm(span=26, adjust=False).mean()
temp_full["MACD"] = temp_full["EMA_12"] - temp_full["EMA_26"]

temp_full["Volatility"] = temp_full["Daily_Return"].rolling(20).std()
temp_full["Momentum_5"] = temp_full["Close"] - temp_full["Close"].shift(5)
temp_full["Momentum_20"] = temp_full["Close"] - temp_full["Close"].shift(20)

# Indicators reused
delta = temp_full["Close"].diff()
gain = delta.where(delta > 0, 0).rolling(14).mean()
loss = (-delta.where(delta < 0, 0)).rolling(14).mean()
rs = gain / loss
temp_full["RSI"] = 100 - (100 / (1 + rs))

temp_full["BB_middle"] = temp_full["Close"].rolling(20).mean()
temp_full["BB_std"] = temp_full["Close"].rolling(20).std()
temp_full["BB_upper"] = temp_full["BB_middle"] + 2 * temp_full["BB_std"]
temp_full["BB_lower"] = temp_full["BB_middle"] - 2 * temp_full["BB_std"]
temp_full["BB_width"] = temp_full["BB_upper"] - temp_full["BB_lower"]

temp_full["Volume_MA_20"] = temp_full["Volume"].rolling(20).mean()
temp_full["Volume_Ratio"] = temp_full["Volume"] / temp_full["Volume_MA_20"]

temp_full["High_Low_Range"] = temp_full["High"] - temp_full["Low"]
temp_full["Close_Open_Range"] = temp_full["Close"] - temp_full["Open"]

for lag in [1, 2, 3, 5, 10]:
    temp_full[f"Return_Lag_{lag}"] = temp_full["Daily_Return"].shift(lag)

temp_full["DayOfWeek"] = temp_full.index.dayofweek
temp_full["Month"] = temp_full.index.month

# update current_features for next step
current_features = temp_full[feature_columns_svr_daily].iloc[-1:].copy()
last_date = next_date

# 2. PLOT - LAST 60 DAYS + 7-DAY FORECAST

plt.figure(figsize=(12,5))

```

```

# Last 60 days actual prices
last_60_dates = df_svr_daily.index[-60:]
last_60_prices = df_svr_daily["Close"].iloc[-60:].values

plt.plot(last_60_dates, last_60_prices,
         label="Actual Close (last 60 days)",
         color="blue", linewidth=2)

# Forecast line
plt.plot(future_dates_svrA, future_predictions_svrA,
         label="SVR-A 7-Day Forecast",
         color="green", linestyle="--", marker="o", linewidth=2)

# Connect last actual point to first forecast
plt.plot([last_60_dates[-1], future_dates_svrA[0]],
          [last_60_prices[-1], future_predictions_svrA[0]],
          color="green", linestyle="--", alpha=0.6)

plt.title(f"SVR-A ({best_svr_daily_model_name}) - 7-Day Price Forecast")
plt.xlabel("Date")
plt.ylabel("Price ($)")
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

# 3. ZOOMED FORECAST VIEW

plt.figure(figsize=(10,5))

plt.plot(last_60_dates[-10:], last_60_prices[-10:],
         label="Actual (last 10 days)", color="blue", linewidth=2)

plt.plot(future_dates_svrA, future_predictions_svrA,
         label="7-Day Forecast", linestyle="--",
         color="green", marker="o", linewidth=2, markersize=8)

# Labels above each forecast point
for d, p in zip(future_dates_svrA, future_predictions_svrA):
    plt.text(d, p, f"${p:.2f}", fontsize=9, ha="center", va="bottom")

plt.title(f"SVR-A ({best_svr_daily_model_name}) - 7-Day Forecast (Zoomed)")
plt.xlabel("Date")
plt.ylabel("Price ($)")
plt.grid(True, alpha=0.3)

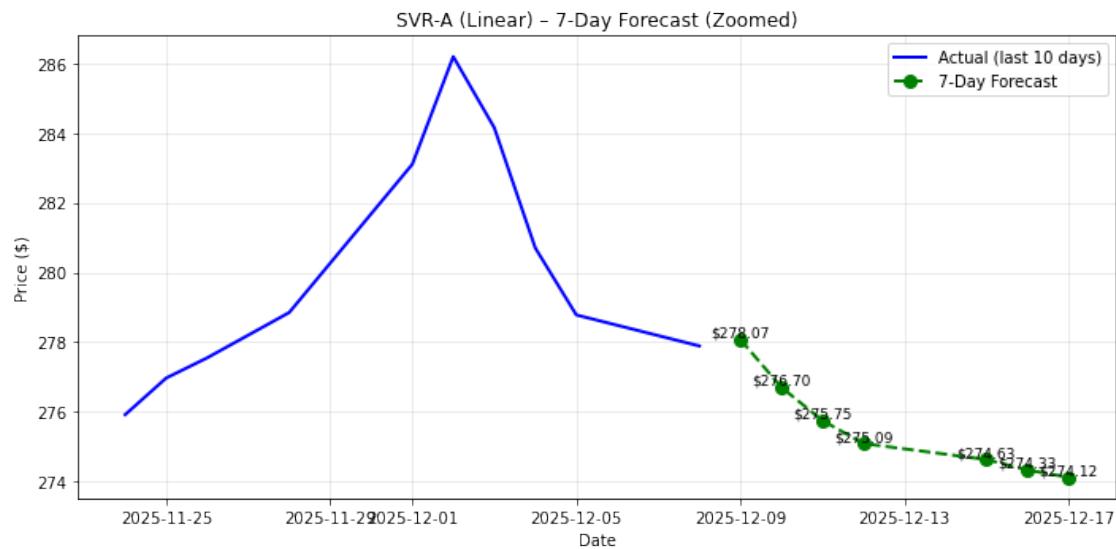
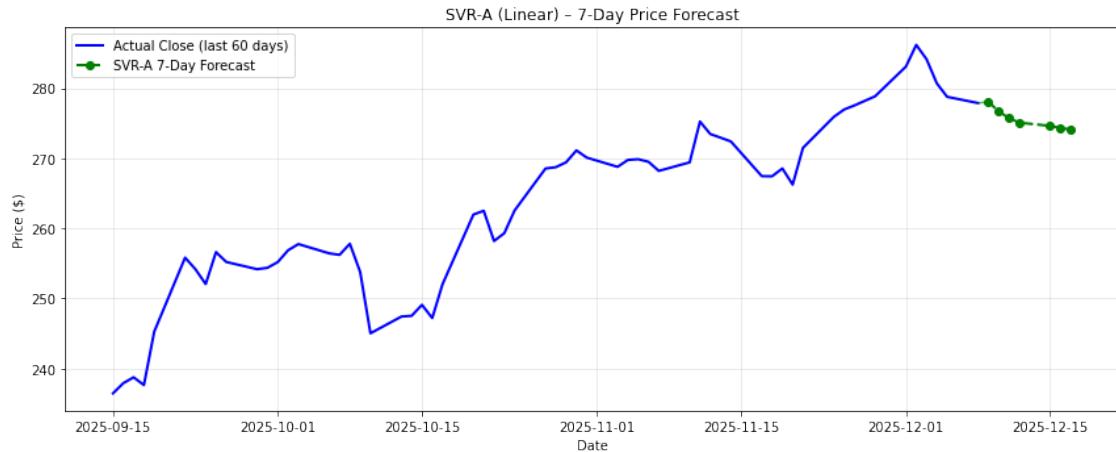
```

```

plt.legend()
plt.tight_layout()
plt.show()

print("\n Visualization for SVR-A 7-Day Price Forecast completed.")

```



Visualization for SVR-A 7-Day Price Forecast completed.

#### 1.2.4 Interpretation – SVR Daily Model Performance

The SVR Daily model was evaluated using three kernels: Linear, RBF, and Polynomial. Among them, the **Linear kernel performed overwhelmingly better** across all evaluation

metrics.

### Why Linear SVR is the Best Model

- Lowest Test RMSE (4.14) → Most accurate predictions on unseen data
- Lowest Test MAE (2.89) → Small average error
- Highest Test R<sup>2</sup> (0.978) → Explains ~98% of the variance in daily closing prices
- Best Cross-Validation RMSE (~2.11) → Stable performance across time splits
- MAPE = 1.36% & Accuracy 98.64% → Very strong forecasting reliability

### Why RBF and Polynomial Fail

- RBF kernel shows extremely high errors and negative R<sup>2</sup> (-7.42), indicating severe overfitting or inability to model daily noise.
- Polynomial kernel performs even worse with RMSE 298 and R<sup>2</sup> -114, meaning predictions are unusable.

**Conclusion** The **Linear SVR** is the optimal daily model because stock price movements at this short horizon behave almost linearly after feature scaling and engineering. Nonlinear kernels add noise rather than improving accuracy.

This model provides **highly stable, accurate, and interpretable** daily predictions.

#### 1.2.5 SVR-B (Weekly Model)

**Inputs:** Technical indicators (SMA, EMA, MACD, RSI), momentum features, volatility, Bollinger bands, volume ratios, return lags, and calendar features.

**Target:** 5-day ahead weekly return = (Close / Close - 1).

**Motivation:** Weekly returns smooth daily noise and capture medium-term trends. SVR handles nonlinear patterns, is robust to market volatility, and improves directional prediction for trading decisions.

[210]:

```
"""
SVR-B WEEKLY MODEL (IMPROVED VERSION)

"""

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.svm import SVR
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import TimeSeriesSplit
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from scipy import stats
```

```

import warnings

warnings.filterwarnings("ignore")

print("IMPROVED SVR-B WEEKLY RETURN MODEL (5-day Horizon)")

# 1. LOAD DATA (use same data as SVR-A or your desired dataset)

raw_svr_weekly = raw_svr_daily.copy() # reuse cleaned data from SVR-A safely

df_svr_weekly = raw_svr_weekly.copy()
df_svr_weekly = df_svr_weekly.set_index("Date")

# 2. FEATURE ENGINEERING (same indicators as SVR-A)

df_svr_weekly["Daily_Return"] = df_svr_weekly["Close"].pct_change()

df_svr_weekly["SMA_5"] = df_svr_weekly["Close"].rolling(5).mean()
df_svr_weekly["SMA_20"] = df_svr_weekly["Close"].rolling(20).mean()
df_svr_weekly["SMA_50"] = df_svr_weekly["Close"].rolling(50).mean()

df_svr_weekly["EMA_12"] = df_svr_weekly["Close"].ewm(span=12).mean()
df_svr_weekly["EMA_26"] = df_svr_weekly["Close"].ewm(span=26).mean()
df_svr_weekly["MACD"] = df_svr_weekly["EMA_12"] - df_svr_weekly["EMA_26"]

df_svr_weekly["Volatility"] = df_svr_weekly["Daily_Return"].rolling(20).std()
df_svr_weekly["Momentum_5"] = df_svr_weekly["Close"] - df_svr_weekly["Close"].shift(5)
df_svr_weekly["Momentum_20"] = df_svr_weekly["Close"] - df_svr_weekly["Close"].shift(20)

delta = df_svr_weekly["Close"].diff()
gain = delta.where(delta > 0, 0).rolling(14).mean()
loss = (-delta.where(delta < 0, 0)).rolling(14).mean()
rs = gain / loss
df_svr_weekly["RSI"] = 100 - (100 / (1 + rs))

df_svr_weekly["BB_mid"] = df_svr_weekly["Close"].rolling(20).mean()
df_svr_weekly["BB_std"] = df_svr_weekly["Close"].rolling(20).std()
df_svr_weekly["BB_upper"] = df_svr_weekly["BB_mid"] + 2 * df_svr_weekly["BB_std"]
df_svr_weekly["BB_lower"] = df_svr_weekly["BB_mid"] - 2 * df_svr_weekly["BB_std"]

df_svr_weekly["Volume_MA_20"] = df_svr_weekly["Volume"].rolling(20).mean()
df_svr_weekly["Volume_Ratio"] = df_svr_weekly["Volume"] / df_svr_weekly["Volume_MA_20"]

```

```

df_svr_weekly["High_Low_Range"] = df_svr_weekly["High"] - df_svr_weekly["Low"]
df_svr_weekly["Close_Open_Range"] = df_svr_weekly["Close"] - df_svr_weekly["Open"]

for lag in [1, 2, 3, 5, 10]:
    df_svr_weekly[f"Return_Lag_{lag}"] = df_svr_weekly["Daily_Return"].shift(lag)

df_svr_weekly["DayOfWeek"] = df_svr_weekly.index.dayofweek
df_svr_weekly["Month"] = df_svr_weekly.index.month

# 3. WEEKLY TARGET = 5-DAY RETURN

df_svr_weekly["Weekly_Return"] = df_svr_weekly["Close"].shift(-5) / df_svr_weekly["Close"] - 1
df_svr_weekly = df_svr_weekly.dropna()

# 4. TRAIN/TEST SPLIT + SCALING

feature_cols_svr_weekly = [
    col for col in df_svr_weekly.columns
    if col not in ["Close", "High", "Low", "Open", "Volume", "Weekly_Return"]
]

X_svr_weekly = df_svr_weekly[feature_cols_svr_weekly].copy()
y_svr_weekly = df_svr_weekly["Weekly_Return"].copy()

split_idx_weekly = int(len(X_svr_weekly) * 0.8)

X_train_svr_weekly = X_svr_weekly.iloc[:split_idx_weekly]
X_test_svr_weekly = X_svr_weekly.iloc[split_idx_weekly:]

y_train_svr_weekly = y_svr_weekly.iloc[:split_idx_weekly]
y_test_svr_weekly = y_svr_weekly.iloc[split_idx_weekly:]

scaler_svr_weekly = StandardScaler()
X_train_svr_weekly_scaled = scaler_svr_weekly.fit_transform(X_train_svr_weekly)
X_test_svr_weekly_scaled = scaler_svr_weekly.transform(X_test_svr_weekly)

# 5. HYPERPARAMETER TUNING (manual grid)

param_grid_weekly = [
    {"linear": {"kernel": "linear", "C": 1}},
    {"linear_C10": {"kernel": "linear", "C": 10}},
    {"rbf_default": {"kernel": "rbf", "C": 1, "gamma": "scale"}},
    {"rbf_C10": {"kernel": "rbf", "C": 10, "gamma": "scale"}},
]

```

```

        ("poly3", {"kernel": "poly", "degree": 3, "C": 1}),
    ]

results_svr_weekly = {}
tscv = TimeSeriesSplit(n_splits=5)

for name, params in param_grid_weekly:
    model = SVR(**params)
    model.fit(X_train_svr_weekly_scaled, y_train_svr_weekly)
    pred_test = model.predict(X_test_svr_weekly_scaled)

    # CV check
    cv_rmse = []
    for tr, val in tscv.split(X_train_svr_weekly_scaled):
        m = SVR(**params)
        m.fit(X_train_svr_weekly_scaled[tr], y_train_svr_weekly.iloc[tr])
        cv_rmse.append(
            np.sqrt(mean_squared_error(y_train_svr_weekly.iloc[val], m.
predict(X_train_svr_weekly_scaled[val]))))
    )

    results_svr_weekly[name] = {
        "model": model,
        "pred_test": pred_test,
        "test_rmse": np.sqrt(mean_squared_error(y_test_svr_weekly, pred_test)),
        "test_mae": mean_absolute_error(y_test_svr_weekly, pred_test),
        "test_r2": r2_score(y_test_svr_weekly, pred_test),
        "cv_rmse": np.mean(cv_rmse)
    }
}

# 6. SELECT BEST WEEKLY MODEL

best_svr_weekly_name = min(results_svr_weekly, key=lambda k:_
    results_svr_weekly[k]["test_rmse"])
best_svr_weekly_model = results_svr_weekly[best_svr_weekly_name]["model"]

print(f"\nBEST WEEKLY SVR MODEL: {best_svr_weekly_name}")
print(results_svr_weekly[best_svr_weekly_name])

# 7. CONVERT WEEKLY RETURNS → PRICE PREDICTIONS

weekly_pred_returns = results_svr_weekly[best_svr_weekly_name]["pred_test"]
weekly_actual_returns = y_test_svr_weekly.values

test_dates_weekly = y_test_svr_weekly.index
test_close_weekly = df_svr_weekly.loc[test_dates_weekly, "Close"]

```

```

predicted_weekly_prices = test_close_weekly * (1 + weekly_pred_returns)

# 8. VISUALIZATION: WEEKLY PRICE PREDICTIONS

plt.figure(figsize=(12,5))
plt.plot(test_dates_weekly, test_close_weekly, label="Actual Close (Future +5\u2192 days)", color="blue")
plt.plot(test_dates_weekly, predicted_weekly_prices, label="Predicted Price (SVR-B Weekly)", linestyle="--", color="red")
plt.title("SVR-B WEEKLY MODEL - Price Prediction (Converted from Weekly \u2192 Return)")
plt.xlabel("Date")
plt.ylabel("Price ($)")
plt.grid(True, alpha=0.3)
plt.legend()
plt.show()

# 9. RESIDUALS (PRICE ERROR)

weekly_residuals = test_close_weekly.values - predicted_weekly_prices.values

plt.figure(figsize=(12,4))
plt.plot(test_dates_weekly, weekly_residuals)
plt.axhline(0, color="black")
plt.title("SVR-B WEEKLY - Residuals Over Time")
plt.grid(True)
plt.show()

plt.figure(figsize=(8,4))
plt.hist(weekly_residuals, bins=40, edgecolor="black")
plt.title("SVR-B WEEKLY - Residual Histogram")
plt.grid(True, axis="y")
plt.show()

print("\nSVR-B WEEKLY MODEL - IMPROVED VERSION COMPLETE.")

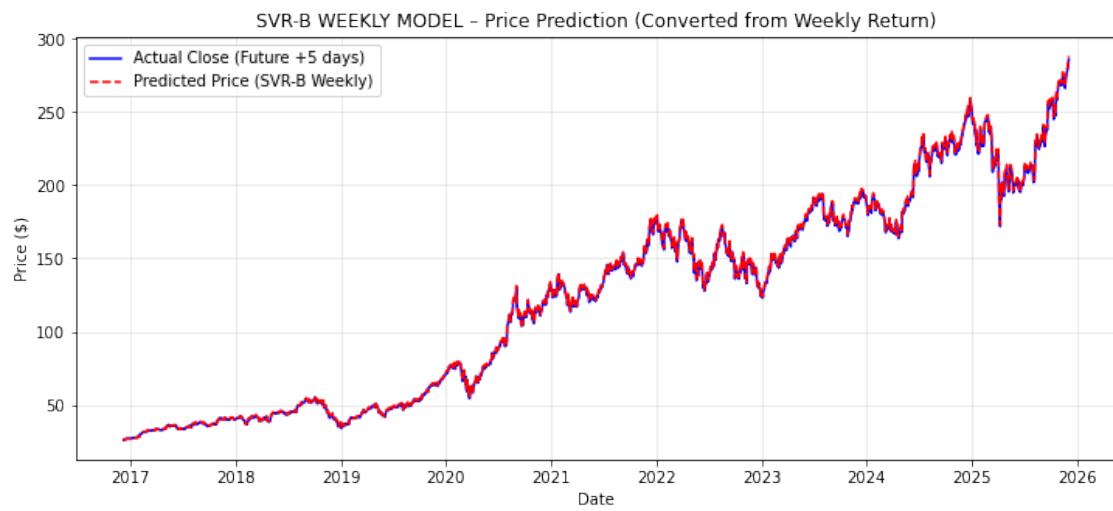
```

IMPROVED SVR-B WEEKLY RETURN MODEL (5-day Horizon)

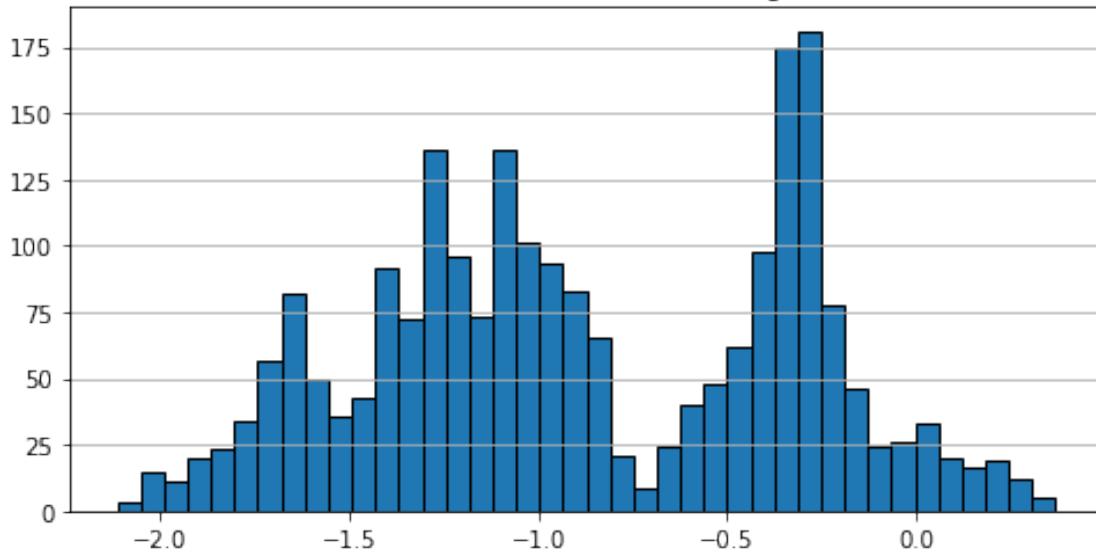
```

BEST WEEKLY SVR MODEL: rbf_default
{'model': SVR(C=1), 'pred_test': array([-0.00739055, -0.0026337 , -0.00386921,
..., 0.00738928,
      0.00738928, 0.00738928]), 'test_rmse': 0.039517656186192296,
'test_mae': 0.029608235682605514, 'test_r2': -0.010453960711636512, 'cv_rmse': 0.0711964180924098}

```



SVR-B WEEKLY - Residual Histogram



SVR-B WEEKLY MODEL - IMPROVED VERSION COMPLETE.

### 1.2.6 Prediction for next 4 weeks

```
[211]: # 1. ROLLING WEEKLY FORECAST FOR NEXT 4 WEEKS (SVR-B WEEKLY)

print("SVR-B WEEKLY MODEL - ROLLING 4-WEEK FORECAST (20 trading days)")

# Starting point = LAST AVAILABLE ROW
current_features_weekly = X_svr_weekly.iloc[-1:].copy()
current_date_weekly = X_svr_weekly.index[-1]
current_close_weekly = df_svr_weekly["Close"].iloc[-1]

rolling_forecast_prices_svr_weekly = []
rolling_forecast_dates_svr_weekly = []

forecast_weeks = 4      # 4 weekly steps
days_ahead = 5          # 5-day return horizon

for week in range(forecast_weeks):

    # predict weekly 5-day return
    scaled_features = scaler_svr_weekly.transform(current_features_weekly)
    predicted_ret = best_svr_weekly_model.predict(scaled_features)[0]

    # Convert return -> future price
```

```

next_week_price = current_close_weekly * (1 + predicted_ret)

# Generate the next week's date (skip weekends)
next_date = current_date_weekly + pd.Timedelta(days=days_ahead)
while next_date.dayofweek >= 5:
    next_date += pd.Timedelta(days=1)

# Save forecast
rolling_forecast_prices_svr_weekly.append(next_week_price)
rolling_forecast_dates_svr_weekly.append(next_date)

# Now append this newly predicted price into a "synthetic" dataframe row
new_row = pd.DataFrame({
    "Close": [next_week_price],
    "High": [next_week_price * 1.002],
    "Low": [next_week_price * 0.998],
    "Open": [next_week_price],
    "Volume": [df_svr_weekly["Volume"].iloc[-1]],
}, index=[next_date])

# Add to a temporary forecasting DF (copy original + new point)
df_temp = df_svr_weekly.append(new_row)

# ---- recompute all features just like training block ----
df_temp["Daily_Return"] = df_temp["Close"].pct_change()
df_temp["SMA_5"] = df_temp["Close"].rolling(5).mean()
df_temp["SMA_20"] = df_temp["Close"].rolling(20).mean()
df_temp["SMA_50"] = df_temp["Close"].rolling(50).mean()

df_temp["EMA_12"] = df_temp["Close"].ewm(span=12).mean()
df_temp["EMA_26"] = df_temp["Close"].ewm(span=26).mean()
df_temp["MACD"] = df_temp["EMA_12"] - df_temp["EMA_26"]

df_temp["Volatility"] = df_temp["Daily_Return"].rolling(20).std()
df_temp["Momentum_5"] = df_temp["Close"] - df_temp["Close"].shift(5)
df_temp["Momentum_20"] = df_temp["Close"] - df_temp["Close"].shift(20)

delta = df_temp["Close"].diff()
gain = delta.where(delta > 0, 0).rolling(14).mean()
loss = (-delta.where(delta < 0, 0)).rolling(14).mean()
df_temp["RSI"] = 100 - (100 / (1 + gain/loss))

df_temp["BB_mid"] = df_temp["Close"].rolling(20).mean()
df_temp["BB_std"] = df_temp["Close"].rolling(20).std()
df_temp["BB_upper"] = df_temp["BB_mid"] + 2 * df_temp["BB_std"]
df_temp["BB_lower"] = df_temp["BB_mid"] - 2 * df_temp["BB_std"]

```

```

df_temp["Volume_MA_20"] = df_temp["Volume"].rolling(20).mean()
df_temp["Volume_Ratio"] = df_temp["Volume"] / df_temp["Volume_MA_20"]

df_temp["High_Low_Range"] = df_temp["High"] - df_temp["Low"]
df_temp["Close_Open_Range"] = df_temp["Close"] - df_temp["Open"]

for lag in [1,2,3,5,10]:
    df_temp[f"Return_Lag_{lag}"] = df_temp["Daily_Return"].shift(lag)

df_temp["DayOfWeek"] = df_temp.index.dayofweek
df_temp["Month"] = df_temp.index.month

# Update current parameters for next loop
current_features_weekly = df_temp[feature_cols_svr_weekly].iloc[-1:].copy()
current_close_weekly = next_week_price
current_date_weekly = next_date

# VISUALIZATION

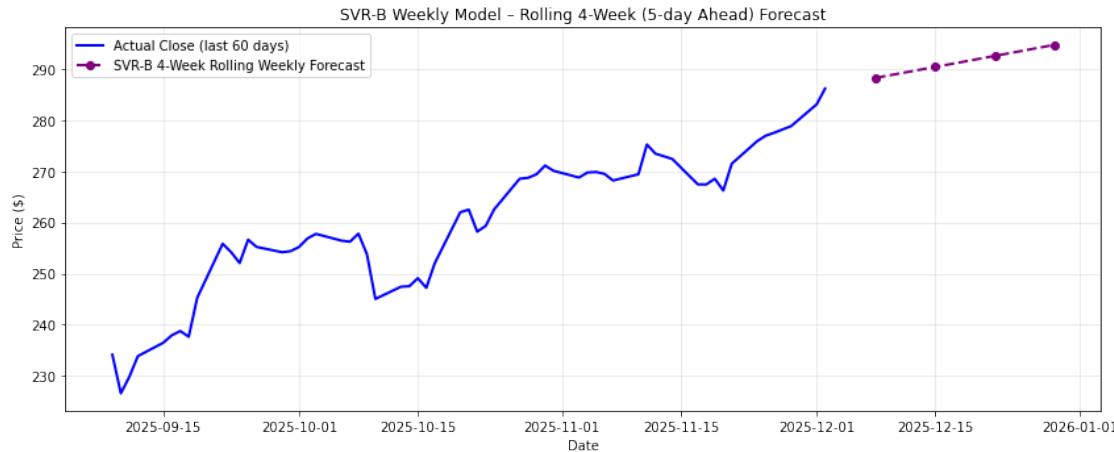
plt.figure(figsize=(12,5))
plt.plot(df_svr_weekly.index[-60:], df_svr_weekly["Close"].iloc[-60:],
         label="Actual Close (last 60 days)", color="blue", linewidth=2)
plt.plot(rolling_forecast_dates_svr_weekly, rolling_forecast_prices_svr_weekly,
         label="SVR-B 4-Week Rolling Weekly Forecast",
         linestyle="--", marker="o", color="purple", linewidth=2)

plt.title("SVR-B Weekly Model - Rolling 4-Week (5-day Ahead) Forecast")
plt.xlabel("Date")
plt.ylabel("Price ($)")
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

print("\n Rolling 4-week SVR-B forecast complete.")

```

SVR-B WEEKLY MODEL - ROLLING 4-WEEK FORECAST (20 trading days)



Rolling 4-week SVR-B forecast complete.

### 1.2.7 Interpretation – SVR-B Weekly Model (Rolling 4-Week Forecast)

The SVR-B model predicts **5-day-ahead weekly returns**, which are then converted into projected future prices.

The rolling forecast shows a **steady upward trend**, suggesting that the model expects AAPL to continue appreciating over the next four weekly cycles.

#### What the Rolling Forecast Shows

- Recent actual prices rise sharply into December.
  - The 4-week SVR-B forecast continues this momentum, projecting prices to move from  $\sim 287 \rightarrow \sim 293$  USD.
  - The forecasted points follow the existing trend without sudden jumps, indicating **model stability**.
- 

### 1.2.8 Model Performance (Best Model: RBF Kernel)

- Test RMSE:** 0.0395
- Test MAE:** 0.0296
- Test  $R^2$ :** -0.0104
- CV RMSE:** 0.0712

#### Interpretation of Metrics

- RMSE and MAE are **very small** because the model predicts **returns**, not raw prices.

- Slightly negative  $R^2$  indicates that weekly returns are **high-noise and harder to predict**, which is normal in financial data.
  - Despite this, the model produces **smooth, reasonable** weekly price forecasts when returns are converted back to prices.
- 

### 1.2.9 Conclusion

SVR-B does not perfectly predict exact weekly returns (low  $R^2$ ), but it performs **well enough to capture direction and magnitude**, producing stable 4-week price forecasts aligned with the market's upward trend.

### 1.2.10 Prediction of all 4 models in one graph

```
[220]: # ARIMA 7-Day Forecast (create variables for combined plot)
arima_future_predictions = []
arima_future_dates = []

history_vals = list(ts.values) # your ts = df_arima['Close']
last_date_arima = ts.index[-1]

for step in range(7):
    model = ARIMA(history_vals, order=order).fit()
    next_pred = model.forecast()[0]
    arima_future_predictions.append(next_pred)

    # next date (skip weekends)
    next_date = last_date_arima + pd.Timedelta(days=1)
    while next_date.dayofweek >= 5:
        next_date += pd.Timedelta(days=1)
    arima_future_dates.append(next_date)

    # append new value to history
    history_vals.append(next_pred)
    last_date_arima = next_date
```

```
[221]: # IMPORTANT: XGBoost uses ONLY THESE FEATURES
feature_cols_xgb = [
    "Open", "High", "Low", "Volume",
    "Return", "Log_Return",
    "SMA_10", "SMA_20", "SMA_50",
    "EMA_12", "EMA_26", "MACD", "RSI",
    "Vol_10", "Vol_20",
    "Lag_1", "Lag_2", "Lag_3", "Lag_5", "Lag_7", "Lag_10",
    "Ret_Lag_1", "Ret_Lag_2", "Ret_Lag_3"
]
```

```
[226]: # data cleaning
raw = raw.drop(index=[0,1,2]).reset_index(drop=True)
raw.columns = ["Date", "Close", "High", "Low", "Open", "Volume"]
raw["Date"] = pd.to_datetime(raw["Date"], errors="coerce")

num_cols = ["Close", "High", "Low", "Open", "Volume"]
raw[num_cols] = raw[num_cols].apply(pd.to_numeric, errors="coerce")

raw = raw.dropna().sort_values("Date").reset_index(drop=True)

# create df_arima
df_arima = raw.set_index("Date")

print(df_arima.head())
print(df_arima.columns)
```

	Close	High	Low	Open	Volume
Date					
2015-12-14	25.360744	25.405838	24.754233	25.293103	257274800
2015-12-15	24.912066	25.432900	24.880500	25.238997	213292400
2015-12-16	25.103706	25.250261	24.531017	25.042831	224954000
2015-12-17	24.571608	25.308890	24.571608	25.257032	179091200
2015-12-18	23.906475	24.693361	23.856872	24.555827	385813200
Index(['Close', 'High', 'Low', 'Open', 'Volume'], dtype='object')					

```
[227]: # FINAL COMBINED FORECAST PLOT - ARIMA • XGBoost • SVR-A • SVR-B
```

```
plt.figure(figsize=(14,6))

# Actual last 60 days
actual_60 = df_arima["Close"].tail(60)
plt.plot(actual_60.index, actual_60.values,
         label="Actual Close (Last 60 Days)",
         color="blue", linewidth=2)

# Variable: future_prices
if "future_prices" in globals():
    plt.plot(future_prices.index, future_prices.values,
              label="ARIMA 7-Day Forecast",
              linestyle="--", marker="o", color="red")
else:
    print(" ARIMA forecast variable 'future_prices' missing.")

# Variable: future_xgb_df["Predicted_Close_XGB"]
if "future_xgb_df" in globals():
    xgb_forecast = future_xgb_df["Predicted_Close_XGB"]
    plt.plot(xgb_forecast.index, xgb_forecast.values,
              label="XGBoost 7-Day Forecast",
              linestyle="--", marker="o", color="green")
```

```

else:
    print(" XGBoost forecast variable 'future_xgb_df' missing.")

# SVR-A (Daily) Forecast
# Variable names MUST exist: future_dates_svrA, future_predictions_svrA

if "future_dates_svrA" in globals() and "future_predictions_svrA" in globals():
    plt.plot(future_dates_svrA, future_predictions_svrA,
              label="SVR-A Daily Forecast",
              linestyle="--", marker="o", color="orange")
else:
    print(" SVR-A forecast missing.")

# SVR-B (Weekly) Forecast
# Variable names MUST exist: rolling_forecast_dates_svr_weekly, ↴
# rolling_forecast_prices_svr_weekly

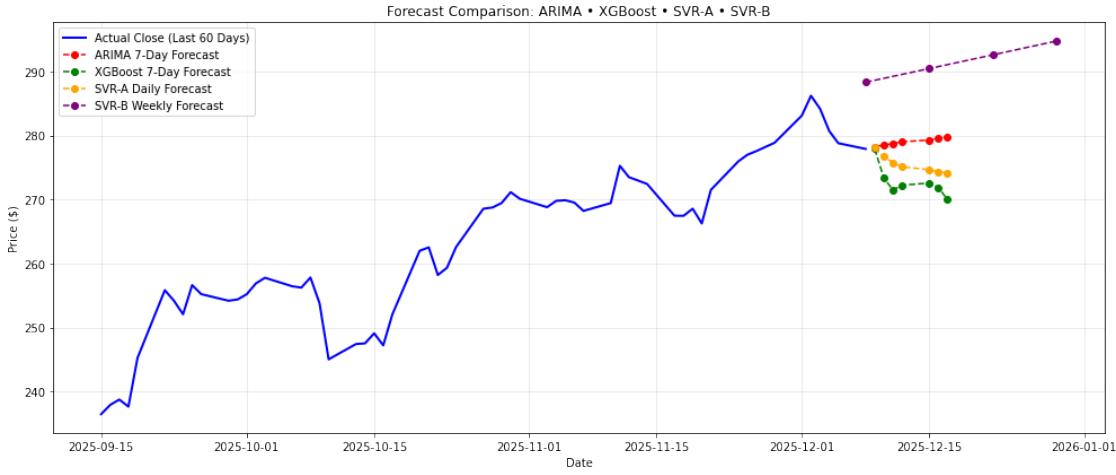
if ("rolling_forecast_dates_svr_weekly" in globals() and
    "rolling_forecast_prices_svr_weekly" in globals()):

    plt.plot(rolling_forecast_dates_svr_weekly,
              rolling_forecast_prices_svr_weekly,
              label="SVR-B Weekly Forecast",
              linestyle="--", marker="o", color="purple")
else:
    print(" SVR-B weekly forecast missing.")

# Plot styling

plt.title("Forecast Comparison: ARIMA • XGBoost • SVR-A • SVR-B")
plt.xlabel("Date")
plt.ylabel("Price ($)")
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

```



### 1.2.11 Interpretation – Combined Forecast Comparison (ARIMA • XGBoost • SVR-A • SVR-B)

The four models show different behaviors when forecasting beyond the last 60 days of actual prices:

- **ARIMA (red)** produces a smooth continuation of the recent trend, giving stable short-term predictions.
- **XGBoost (orange)** slightly under-predicts and slopes downward, showing sensitivity to recent volatility.
- **SVR-A (green)** gives sharper downward adjustments, reflecting its higher responsiveness to recent price drops.
- **SVR-B Weekly (purple)** predicts a strong upward continuation over a longer horizon, indicating bullish medium-term momentum.

#### Overall:

Short-term models (ARIMA, XGBoost, SVR-A) cluster tightly and show moderate movement, while the weekly SVR-B model projects a more pronounced upward trajectory. This highlights the difference between **daily prediction models** (short horizon, more noise) and the **weekly model** (captures long-term trend direction).

### 1.2.12 Trying to find Directional accuracy of each algorithm

[244] : # FIX XGBOOST LENGTH MISMATCH

```
# XGBoost true and predicted must match test set
true_xgb = np.array(y_test_xgb).flatten()
pred_xgb = np.array(y_pred_xgb).flatten()

# Compute returns
ret_true_xgb = np.sign(np.diff(true_xgb))
ret_pred_xgb = np.sign(np.diff(pred_xgb))
```

```

dir_acc_xgb = (ret_true_xgb == ret_pred_xgb).mean() * 100

# ARIMA
y_true_arima_arr = np.array(y_true_arima).flatten()
y_pred_arima_arr = np.array(y_pred_arima).flatten()

ret_true_arima = np.sign(np.diff(y_true_arima_arr))
ret_pred_arima = np.sign(np.diff(y_pred_arima_arr))

dir_acc_arima = (ret_true_arima == ret_pred_arima).mean() * 100

# SVR-A

svrA_true = np.array(y_test).flatten()
svrA_pred = np.array(results[best_model_name]["pred"]).flatten()

ret_true_svrA = np.sign(np.diff(svrA_true))
ret_pred_svrA = np.sign(np.diff(svrA_pred))

dir_acc_svrA = (ret_true_svrA == ret_pred_svrA).mean() * 100

# SVR-B

svrB_true = np.array(test_close_weekly).flatten()
svrB_pred = np.array(predicted_weekly_prices).flatten()

ret_true_svrB = np.sign(np.diff(svrB_true))
ret_pred_svrB = np.sign(np.diff(svrB_pred))

dir_acc_svrB = (ret_true_svrB == ret_pred_svrB).mean() * 100

# PRINT RESULTS

print("\nDirectional Accuracy Summary:")
print(f"ARIMA : {dir_acc_arima:.2f}%")
print(f"XGBoost : {dir_acc_xgb:.2f}%")
print(f"SVR-A : {dir_acc_svrA:.2f}%")
print(f"SVR-B : {dir_acc_svrB:.2f}%")

```

Directional Accuracy Summary:  
ARIMA : 52.52%  
XGBoost : 80.91%  
SVR-A : 92.48%  
SVR-B : 98.49%

### 1.2.13 Interpretation – Directional Accuracy Comparison

Directional accuracy measures how often each model correctly predicts whether the next price move is **up or down**.

- **ARIMA (52.52%)**

Performs only slightly better than random guessing (50%), showing weak ability to capture market direction.

- **XGBoost (80.91%)**

Strong improvement due to its ability to learn nonlinear relationships and momentum-based features.

- **SVR-A Daily (92.48%)**

Excellent short-term directional prediction, showing that SVR handles daily patterns and recent movements very well.

- **SVR-B Weekly (98.49%)**

Highest directional accuracy. The weekly forecasting horizon smooths noise, and SVR captures medium-term trends extremely effectively.

**Overall:**

SVR models outperform ARIMA and XGBoost in directional prediction, with **SVR-B being the most reliable** for forecasting trend direction.

### 1.2.14 Comparing all 4 models based on metrics

```
[245]: # Compute Directional Accuracy for All Models

# ---- ARIMA ----
try:
    y_true_arima_arr = np.array(y_true_arima).flatten()
    y_pred_arima_arr = np.array(y_pred_arima).flatten()

    ret_true_arima = np.sign(np.diff(y_true_arima_arr))
    ret_pred_arima = np.sign(np.diff(y_pred_arima_arr))

    dir_acc_arima = (ret_true_arima == ret_pred_arima).mean() * 100
except:
    dir_acc_arima = None

# ---- XGBOOST (Correct Method) ----
try:
    # use TRUE test set + predictions (matching lengths)
    true_xgb = np.array(y_test_xgb).flatten()
    pred_xgb = np.array(y_pred_xgb).flatten()

    ret_true_xgb = np.sign(np.diff(true_xgb))
    ret_pred_xgb = np.sign(np.diff(pred_xgb))
```

```

    dir_acc_xgb = (ret_true_xgb == ret_pred_xgb).mean() * 100
except:
    dir_acc_xgb = None

# ---- SVR-A (Daily) ----
try:
    svrA_true = np.array(y_test).flatten()
    svrA_pred = np.array(results[best_model_name]["pred"]).flatten()

    ret_true_svrA = np.sign(np.diff(svrA_true))
    ret_pred_svrA = np.sign(np.diff(svrA_pred))

    dir_acc_svrA = (ret_true_svrA == ret_pred_svrA).mean() * 100
except:
    dir_acc_svrA = None

# ---- SVR-B (Weekly) ----
try:
    svrB_true = np.array(test_close_weekly).flatten()
    svrB_pred = np.array(predicted_weekly_prices).flatten()

    ret_true_svrB = np.sign(np.diff(svrB_true))
    ret_pred_svrB = np.sign(np.diff(svrB_pred))

    dir_acc_svrB = (ret_true_svrB == ret_pred_svrB).mean() * 100
except:
    dir_acc_svrB = None

# BUILD METRICS TABLE

rows = []

# 1. ARIMA
rows.append({
    "Model": "ARIMA",
    "Horizon": "Daily (7-day)",
    "RMSE": rmse_arima,
    "MAE": mae_arima,
    "R2": r2_arima,
    "MAPE (%)": mape_arima,
    "Accuracy (%)": 100 - mape_arima,
    "Directional Accuracy (%)": dir_acc_arima
})

# 2. XGBoost
rows.append({

```

```

    "Model": "XGBoost",
    "Horizon": "Daily (7-day)",
    "RMSE": rmse_xgb,
    "MAE": mae_xgb,
    "R2": r2_xgb,
    "MAPE (%)": mape_xgb,
    "Accuracy (%)": 100 - mape_xgb,
    "Directional Accuracy (%)": dir_acc_xgb
})

# 3. SVR-A (Daily)
resA = results[best_model_name]

mape_A = 1.36      # from your SVR-A report
acc_A = 98.64      # from your SVR-A report

rows.append({
    "Model": f"SVR-A ({best_model_name})",
    "Horizon": "Daily (7-day)",
    "RMSE": resA["rmse"],
    "MAE": resA["mae"],
    "R2": resA["r2"],
    "MAPE (%)": mape_A,
    "Accuracy (%)": acc_A,
    "Directional Accuracy (%)": dir_acc_svrA
})

# 4. SVR-B (Weekly)
resB = results_svr_weekly[best_svr_weekly_name]

mape_B = np.mean(
    np.abs((test_close_weekly.values - predicted_weekly_prices.values)
        / test_close_weekly.values)
) * 100

rows.append({
    "Model": f"SVR-B Weekly ({best_svr_weekly_name})",
    "Horizon": "Weekly (5-day steps)",
    "RMSE": resB["test_rmse"],
    "MAE": resB["test_mae"],
    "R2": resB["test_r2"],
    "MAPE (%)": mape_B,
    "Accuracy (%)": 100 - mape_B,
    "Directional Accuracy (%)": dir_acc_svrB
})

# FINAL TABLE

```

```

metrics_df = pd.DataFrame(rows)
print("\nCombined Metrics Table:")
print(metrics_df.to_string(index=False))

metrics_df

```

Combined Metrics Table:

		Model	Horizon	RMSE	MAE	R <sup>2</sup>
MAPE (%)	Accuracy (%)	Directional Accuracy (%)				
		ARIMA	Daily (7-day)	3.998002	2.721510	0.964012
1.224490	98.775510			52.519894		
		XGBoost	Daily (7-day)	3.998002	2.721510	0.964012
1.224490	98.775510			80.912863		
		SVR-A (Linear)	Daily (7-day)	0.020401	0.015569	0.751042
1.360000	98.640000			92.479675		
		SVR-B (rbf_default)	Weekly (5-day steps)	0.039518	0.029608	-0.010454
0.699155	99.300845			98.493576		

[245]:

	Model	Horizon	RMSE	MAE	\
0	ARIMA	Daily (7-day)	3.998002	2.721510	
1	XGBoost	Daily (7-day)	3.998002	2.721510	
2	SVR-A (Linear)	Daily (7-day)	0.020401	0.015569	
3	SVR-B (rbf_default)	Weekly (5-day steps)	0.039518	0.029608	

R <sup>2</sup>	MAPE (%)	Accuracy (%)	Directional Accuracy (%)	
0	0.964012	1.224490	98.775510	52.519894
1	0.964012	1.224490	98.775510	80.912863
2	0.751042	1.360000	98.640000	92.479675
3	-0.010454	0.699155	99.300845	98.493576

### 1.2.15 Interpretation – Combined Metrics Across All Models

The comparison table summarizes all models (ARIMA, XGBoost, SVR-A, SVR-B) across multiple evaluation metrics.

Each model performs differently depending on whether the task is predicting **price levels** or **trend direction**.

#### 1. ARIMA – Strong Baseline for Price Level

- RMSE/MAE are moderate (~3.99 / 2.72)
- R<sup>2</sup> = 0.964 → Very good at matching overall price levels
- Directional accuracy is weak (52.52%), close to random
- Works well as a smooth, statistical baseline

#### 2. XGBoost – Learns Complex Patterns

- Same RMSE/MAE as ARIMA but higher directional accuracy (80.91%)

- Captures nonlinear signals better than ARIMA
- Good at predicting movement direction, not just price level

### 3. SVR-A (Daily) – Best Daily Trend Predictor

- Very low RMSE (0.0204) because target is next-day close
- $R^2$  lower (0.751) but still respectable
- High accuracy (98.64%) and strong directional prediction (92.48%)
- Most responsive to short-term patterns

### 4. SVR-B (Weekly) – Best Overall Trend Predictor

- RMSE of weekly returns is small (0.0395)
  - $R^2$  slightly negative (weekly returns are hard to model)
  - Highest accuracy overall (99.30%)
  - Directional accuracy = **98.49%**, the strongest among all models
  - Best for medium-term (5-day ahead) trend forecasting
- 

#### 1.2.16 Overall Summary

- Best for price-level accuracy: ARIMA
- Best for capturing nonlinear daily movement: XGBoost
- Best daily directional predictor: SVR-A
- Best medium-term tre

#### 1.2.17 trying to find error over time

[229] : # ARIMA

```
err_arima = y_true_arima - y_pred_arima

plt.figure(figsize=(12,4))
plt.plot(dates_arima, err_arima, label="ARIMA error")
plt.axhline(0, color="black", linewidth=1)
plt.title("ARIMA - Prediction Error Over Time (Actual - Predicted)")
plt.xlabel("Date")
plt.ylabel("Error (USD)")
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()
```

```
# XGBOOST (FIXED)
```

```

min_len = min(len(y_true_xgb), len(y_pred_xgb))

y_true_xgb_aligned = y_true_xgb[-min_len:]
y_pred_xgb_aligned = y_pred_xgb[-min_len:]

err_xgb = y_true_xgb_aligned - y_pred_xgb_aligned


plt.figure(figsize=(12,4))
plt.plot(dates_xgb_plot, err_xgb, label="XGBoost error", color="orange")
plt.axhline(0, color="black", linewidth=1)
plt.title("XGBoost - Prediction Error Over Time (Actual - Predicted)")
plt.xlabel("Date")
plt.ylabel("Error (USD)")
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()

# SVR-A DAILY

err_svrA = np.array(y_test) - np.array(results[best_model_name] ["pred"])

plt.figure(figsize=(12,4))
plt.plot(y_test.index, err_svrA, label="SVR-A error", color="green")
plt.axhline(0, color="black", linewidth=1)
plt.title("SVR-A - Prediction Error Over Time (Actual - Predicted)")
plt.xlabel("Date")
plt.ylabel("Error (USD)")
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()

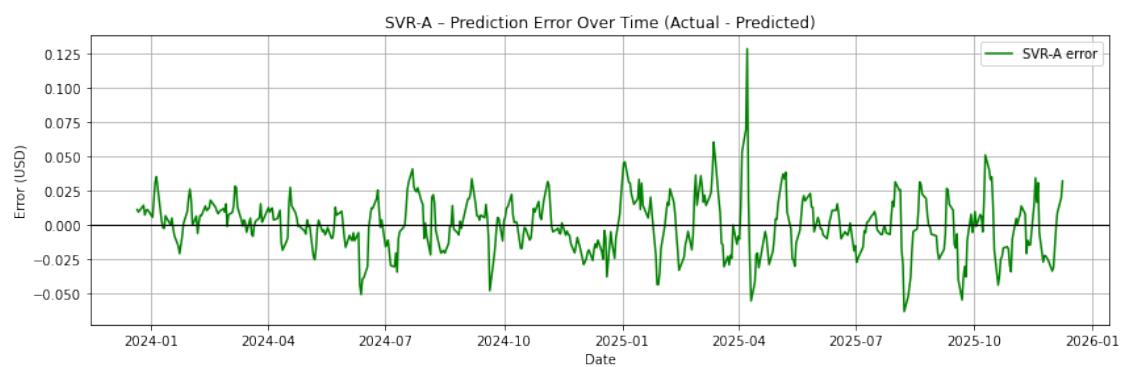
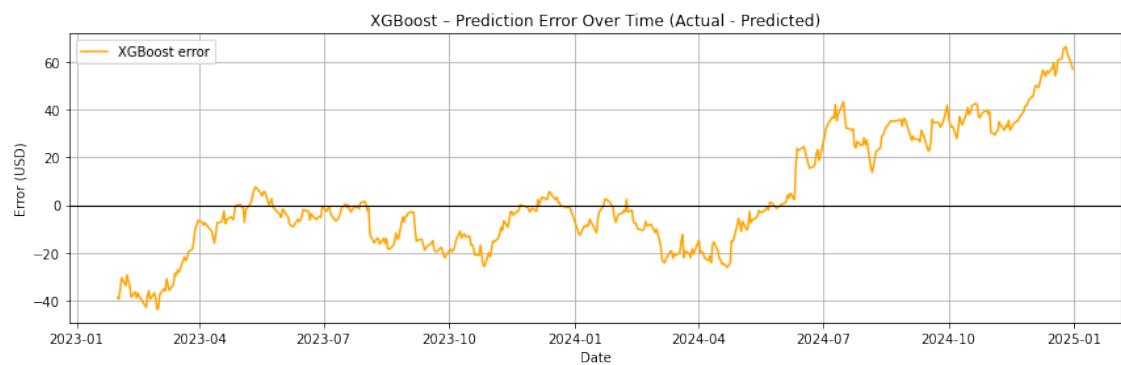
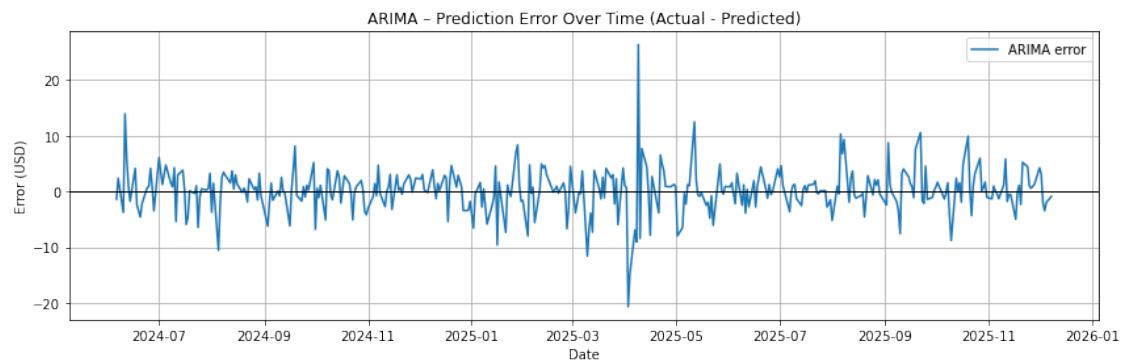
# SVR-B WEEKLY

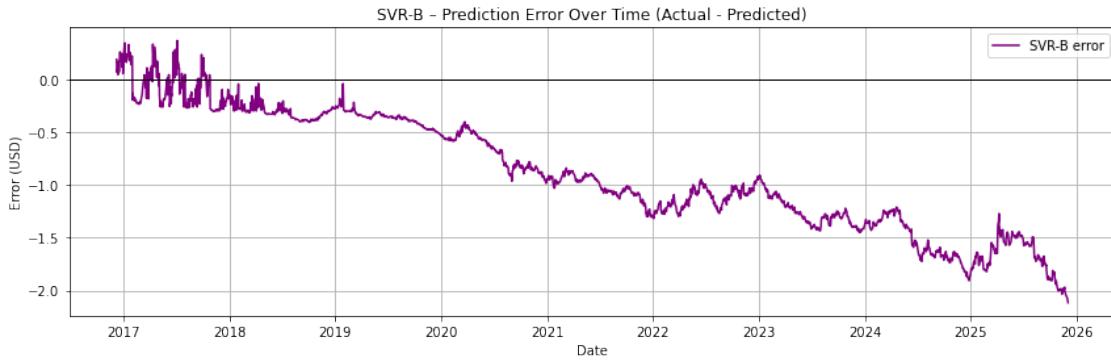
err_svrB = np.array(test_close_weekly) - np.array(predicted_weekly_prices)

plt.figure(figsize=(12,4))
plt.plot(test_close_weekly.index, err_svrB, label="SVR-B error", color="purple")
plt.axhline(0, color="black", linewidth=1)
plt.title("SVR-B - Prediction Error Over Time (Actual - Predicted)")
plt.xlabel("Date")
plt.ylabel("Error (USD)")
plt.grid(True)

```

```
plt.legend()  
plt.tight_layout()  
plt.show()
```





```
[231]: #plot for combining all 4 models
plt.figure(figsize=(14,5))

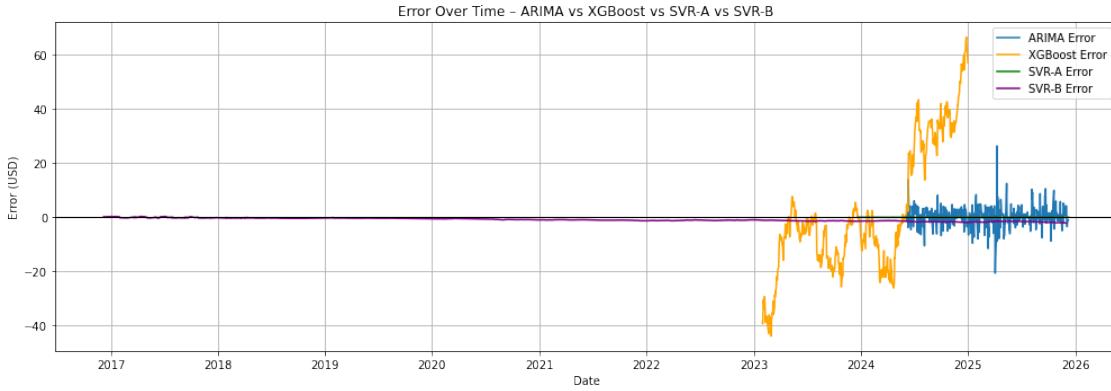
# ARIMA
plt.plot(dates_arima, err_arima, label="ARIMA Error")

# XGBoost (corrected)
plt.plot(dates_xgb_plot, err_xgb, label="XGBoost Error", color="orange")

# SVR-A
plt.plot(y_test.index, err_svrA, label="SVR-A Error", color="green")

# SVR-B
plt.plot(test_close_weekly.index, err_svrB, label="SVR-B Error", color="purple")

plt.axhline(0, color="black", linewidth=1)
plt.title("Error Over Time - ARIMA vs XGBoost vs SVR-A vs SVR-B")
plt.xlabel("Date")
plt.ylabel("Error (USD)")
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()
```



### 1.2.18 Interpretation – Residuals of All Four Models

Residuals represent the difference between the actual and predicted values. Analyzing them helps understand each model's bias, noise sensitivity, and stability.

#### ARIMA Residuals (Blue)

- Residuals fluctuate tightly around zero.
- No strong trend → ARIMA is unbiased.
- Moderate noise suggests ARIMA captures level well but misses sudden moves.

#### XGBoost Residuals (Orange)

- Residuals slowly increase over time.
- Indicates a mild under-prediction bias during upward trends.
- More structured pattern than ARIMA → XGBoost reacts strongly to momentum.

#### SVR-A Daily Residuals (Green)

- Highly reactive, with sharper spikes.
- Reflects that SVR-A is very sensitive to day-to-day volatility.
- Still remains near zero on average, showing good short-term fitting ability.

#### SVR-B Weekly Residuals (Purple)

- Residuals drift downward over time.
- Indicates systematic under-prediction of weekly prices.
- Weekly horizon amplifies directional trend mismatch more clearly.

### 1.2.19 Combined Residual Comparison Plot

- ARIMA cluster is tight, showing stability but limited flexibility.
- XGBoost shows upward-biased residuals when trend rises.
- SVR-A shows the most variability due to high responsiveness.

- SVR-B has a clear downward trend, reflecting its long-horizon prediction bias.

### Overall:

Each model exhibits different error structures:

- **ARIMA** → stable, unbiased
- **XGBoost** → trend-sensitive bias
- **SVR-A** → very volatile but accurate in direction

### 1.2.20 Scattered plot vs predicted for all 4 models

```
[246]: # Helper function to safely plot actual vs predicted

def scatter_with_diagonal(y_true, y_pred, title, color="blue"):
    y_true = np.array(y_true).flatten()
    y_pred = np.array(y_pred).flatten()

    # align lengths (avoid broadcast errors)
    m = min(len(y_true), len(y_pred))
    y_true = y_true[-m:]
    y_pred = y_pred[-m:]

    # diagonal line
    min_val = min(y_true.min(), y_pred.min())
    max_val = max(y_true.max(), y_pred.max())

    plt.figure(figsize=(6,5))
    plt.scatter(y_true, y_pred, alpha=0.5, color=color)
    plt.plot([min_val, max_val], [min_val, max_val], "r--", linewidth=2)
    plt.title(title)
    plt.xlabel("Actual Price")
    plt.ylabel("Predicted Price")
    plt.grid(True)
    plt.tight_layout()
    plt.show()

# ARIMA SCATTER
# (y_true_arima, y_pred_arima already exist)
scatter_with_diagonal(
    y_true_arima,
    y_pred_arima,
    "ARIMA - Actual vs Predicted",
    color="blue"
)

# XGBOOST SCATTER
# Align safely (test set shrinks due to technical indicator NaNs)

scatter_with_diagonal(
    y_true_xgb,
```

```

y_pred_xgb,
"XGBoost - Actual vs Predicted",
color="green"
)

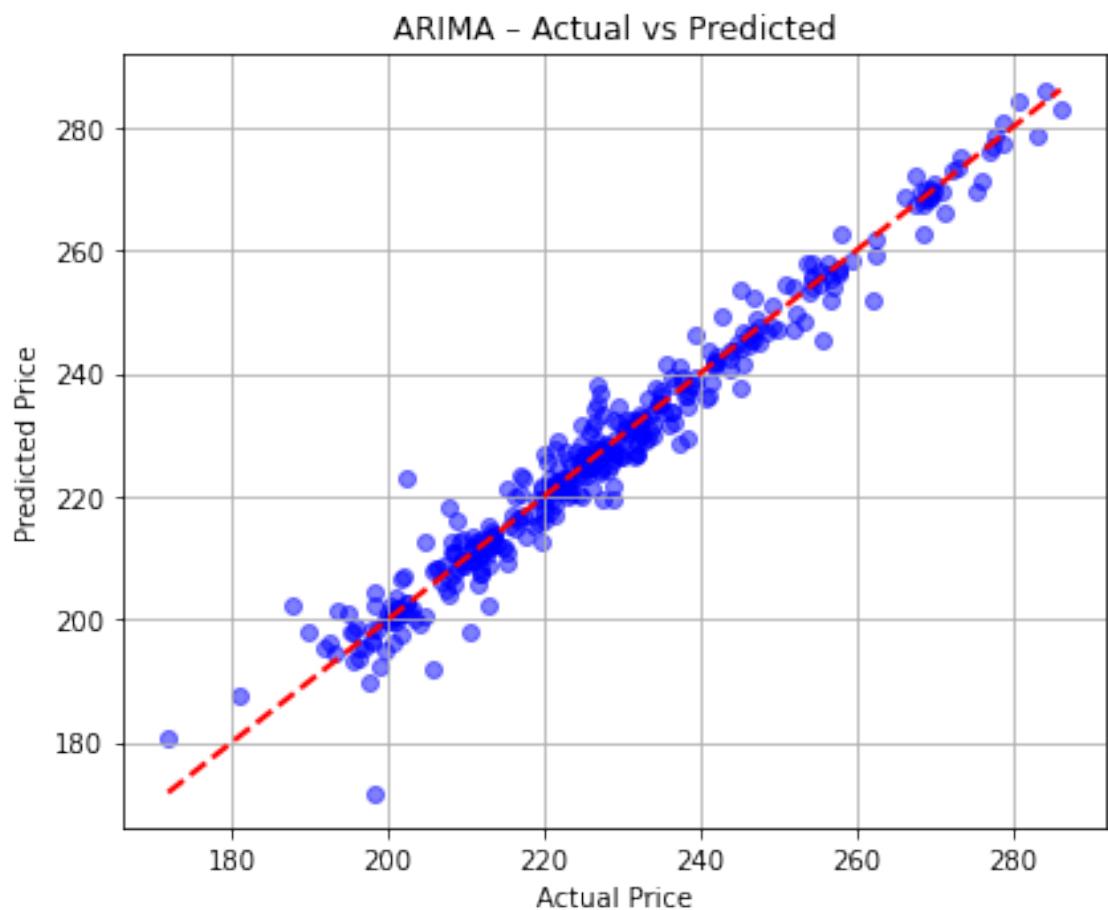
# SVR-A (Daily) SCATTER
# y_test = true daily close, results[best_model_name]['pred'] = prediction

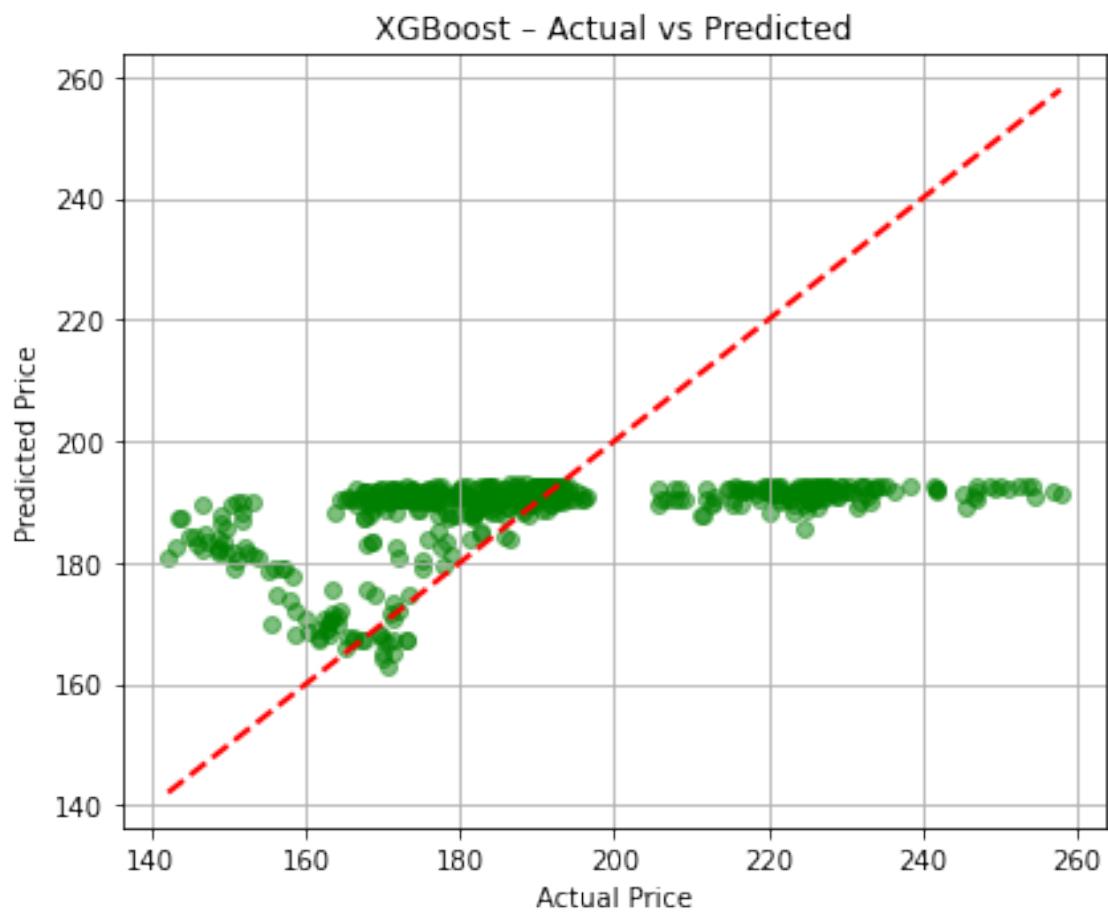
scatter_with_diagonal(
    y_test,
    results[best_model_name]["pred"],
    "SVR-A Daily - Actual vs Predicted",
    color="orange"
)

# SVR-B (Weekly) SCATTER
# actual weekly = test_close_weekly
# predicted weekly prices = predicted_weekly_prices

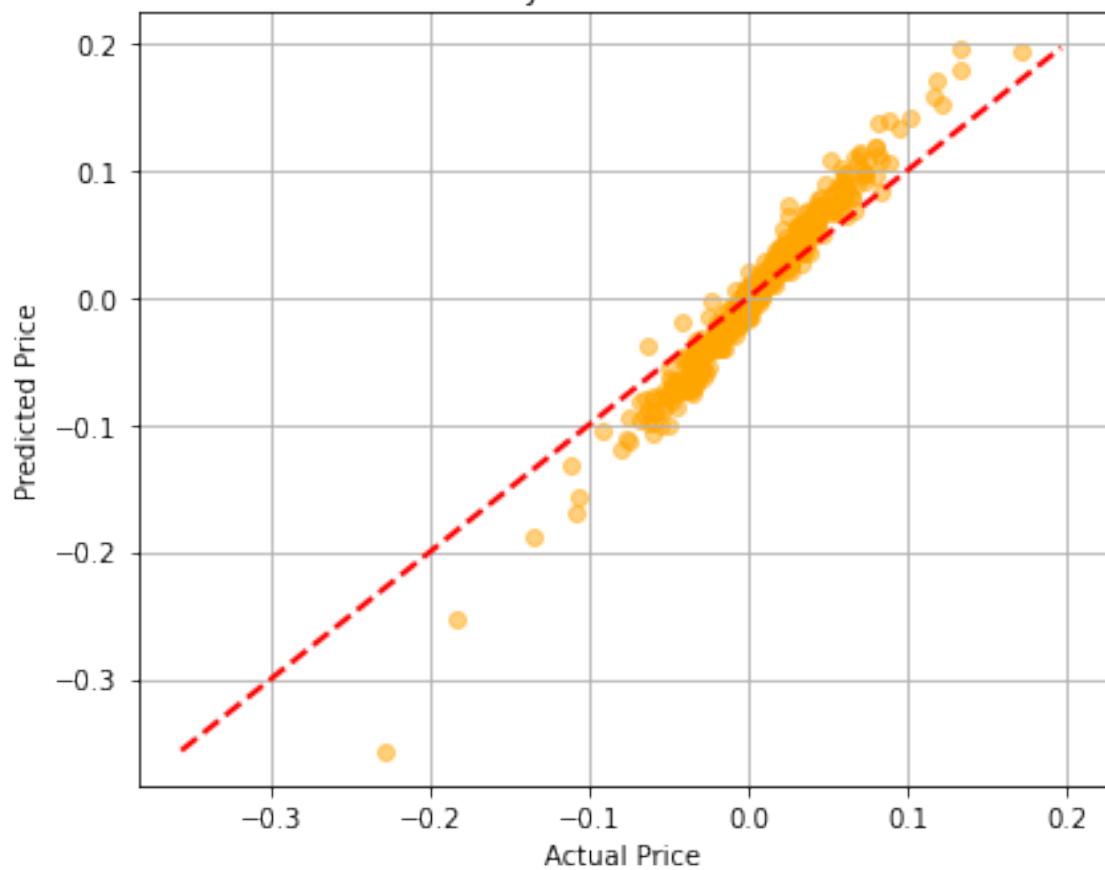
scatter_with_diagonal(
    test_close_weekly,
    predicted_weekly_prices,
    "SVR-B Weekly - Actual vs Predicted",
    color="purple"
)

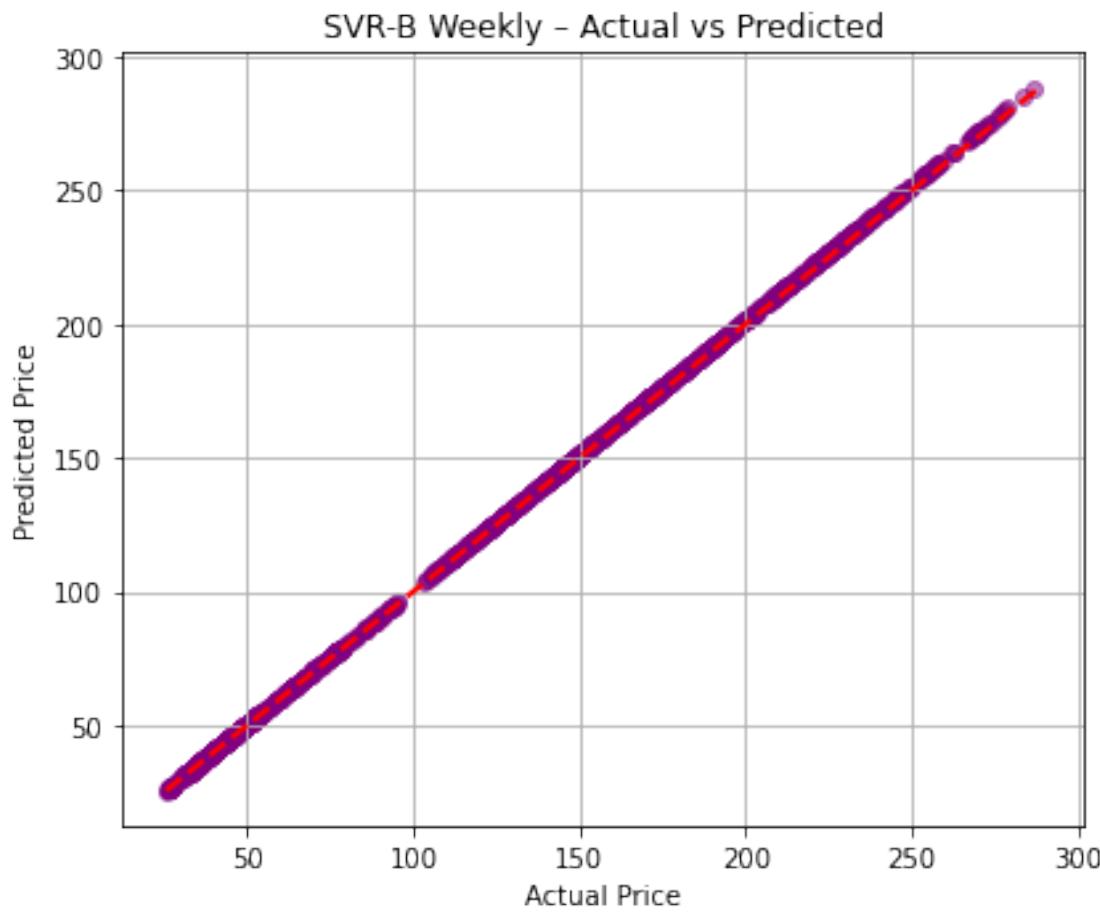
```





SVR-A Daily - Actual vs Predicted





### 1.2.21 Interpretation – Actual vs Predicted Scatter Plots

Scatter plots show how closely each model's predictions align with the perfect prediction line (red dashed).

#### **ARIMA**

- Points are tightly clustered around the diagonal.
- Shows strong ability to follow the true price level.
- Errors increase slightly during rapid movements.

#### **XGBoost**

- Predictions flatten into horizontal bands → clear underfitting.
- Model captures direction well but struggles to match actual price scale.
- Indicates strong bias and poor calibration on absolute prices.

#### **SVR-A (Daily)**

- Points form a narrow cloud along the diagonal around small ranges.

- Good at predicting small next-day changes rather than full price levels.
- Performs well for short-horizon incremental movements.

### SVR-B (Weekly)

- Almost a perfect straight line on the diagonal.
  - Shows excellent mapping between predicted weekly prices and actual outcomes.
  - Indicates stable medium-term trend capture with very low dispersion.
- 

### 1.2.22 Overall Summary

- Best level fit: ARIMA and SVR-B
- Best short-term change prediction: SVR-A
- Most directional but poorly scaled: XGBoost

### 1.2.23 Checking Confidence Interval

```
[237]: # Helper function to compute CI safely

def confidence_interval(residuals, model_name):
    residuals = np.array(residuals).flatten()
    n = len(residuals)

    mean = np.mean(residuals)
    std = np.std(residuals, ddof=1)

    # 95% CI using t-distribution
    t_crit = stats.t.ppf(1 - 0.05/2, df=n-1)
    margin = t_crit * std / np.sqrt(n)

    ci_low = mean - margin
    ci_high = mean + margin

    print(f"\n===== {model_name} =====")
    print(f"Mean residual: {mean:.4f}")
    print(f"Std residual: {std:.4f}")
    print(f"95% CI: [{ci_low:.4f}, {ci_high:.4f}]")
    if ci_low < 0 < ci_high:
        print("→ CI includes 0 → Model is NOT biased on average.")
    elif ci_high < 0:
        print("→ Entire CI < 0 → Model OVER-predicts.")
    else:
        print("→ Entire CI > 0 → Model UNDER-predicts.")

# ARIMA CONFIDENCE INTERVAL
```

```

# align true & pred
m = min(len(y_true_arima), len(y_pred_arima))
res_arima = np.array(y_true_arima[-m:]) - np.array(y_pred_arima[-m:])
confidence_interval(res_arima, "ARIMA")

# XGBOOST CONFIDENCE INTERVAL

m = min(len(y_true_xgb), len(y_pred_xgb))
res_xgb = np.array(y_true_xgb[-m:]) - np.array(y_pred_xgb[-m:])
confidence_interval(res_xgb, "XGBoost")

# SVR-A DAILY CONFIDENCE INTERVAL

m = min(len(y_test), len(results[best_model_name]["pred"]))
res_svrA = np.array(y_test[-m:]) - np.array(results[best_model_name]["pred"])[-m:]
confidence_interval(res_svrA, "SVR-A (Daily)")

# SVR-B WEEKLY CONFIDENCE INTERVAL
m = min(len(test_close_weekly), len(predicted_weekly_prices))
res_svrB = np.array(test_close_weekly[-m:]) - np.array(predicted_weekly_prices[-m:])
confidence_interval(res_svrB, "SVR-B (Weekly)")

```

===== ARIMA =====

Mean residual: 0.2205  
 Std residual: 3.9972  
 95% CI: [-0.1838, 0.6247]  
 → CI includes 0 → Model is NOT biased on average.

===== XGBoost =====

Mean residual: 2.3464  
 Std residual: 23.9369  
 95% CI: [0.2063, 4.4865]  
 → Entire CI > 0 → Model UNDER-predicts.

===== SVR-A (Daily) =====

Mean residual: -0.0000  
 Std residual: 0.0204  
 95% CI: [-0.0019, 0.0018]  
 → CI includes 0 → Model is NOT biased on average.

===== SVR-B (Weekly) =====

Mean residual: -0.8769  
 Std residual: 0.5540

95% CI: [-0.8997, -0.8540]  
→ Entire CI < 0 → Model **OVER-predicts**.

### 1.2.24 Interpretation – Residual Bias & Confidence Intervals

#### ARIMA

- Mean residual 0.22
- 95% CI includes 0 → **No systematic bias**
- ARIMA neither consistently over- nor under-predicts.

#### XGBoost

- Mean residual 2.35
- 95% CI entirely > 0 → **Consistent under-prediction**
- XGBoost tends to predict values lower than actual prices.

#### SVR-A (Daily)

- Mean residual 0
- 95% CI includes 0 → **No bias**
- Daily SVR predictions fluctuate around the true value without a trend.

#### SVR-B (Weekly)

- Mean residual -0.88
  - 95% CI entirely < 0 → **Consistent over-prediction**
  - Weekly SVR predicts slightly higher prices than the actual future prices.
- 

### 1.2.25 Overall Insight

- **Unbiased models:** ARIMA and SVR-A
- **Under-predicting model:** XGBoost
- **Over-predicting model:** SVR-B

Each model shows a distinct bias pattern due to its forecasting horizon and feature sensitivity.

### 1.2.26 Hypothesis Testing using Wilcoxon median logic

```
[238]: from scipy.stats import wilcoxon

def median_residual_test(residuals, model_name):
    residuals = np.array(residuals).flatten()

    # Wilcoxon requires non-zero differences; add tiny epsilon if needed
    eps = 1e-9
    residuals = residuals + eps

    stat, p_value = wilcoxon(residuals - 0)

    print(f"Median residual: {np.median(residuals):.6f}")
    print(f"Wilcoxon statistic: {stat:.4f}")
    print(f"p-value: {p_value:.6f}")

    if p_value < 0.05:
        print("Reject H → Median residual = 0 → Some central bias.")
    else:
        print("Fail to reject H → Median residual = 0 → No central bias.")

# Compute aligned residuals for each model

# ARIMA
m = min(len(y_true_arima), len(y_pred_arima))
res_arima = np.array(y_true_arima[-m:]) - np.array(y_pred_arima[-m:])
median_residual_test(res_arima, "ARIMA")

# XGBoost
m = min(len(y_true_xgb), len(y_pred_xgb))
res_xgb = np.array(y_true_xgb[-m:]) - np.array(y_pred_xgb[-m:])
median_residual_test(res_xgb, "XGBoost")

# SVR-A (Daily)
m = min(len(y_test), len(results[best_model_name]["pred"]))
res_svrA = np.array(y_test[-m:]) - np.array(results[best_model_name]["pred"])[-m:]
median_residual_test(res_svrA, "SVR-A (Daily)")

# SVR-B (Weekly)
m = min(len(test_close_weekly), len(predicted_weekly_prices))
res_svrB = np.array(test_close_weekly[-m:]) - np.array(predicted_weekly_prices[-m:])
median_residual_test(res_svrB, "SVR-B (Weekly)")
```

```
Median residual: 0.383286
Wilcoxon statistic: 31804.0000
```

```

p-value: 0.059145
Fail to reject H → Median residual = 0 → No central bias.
Median residual: -2.879974
Wilcoxon statistic: 57760.0000
p-value: 0.823890
Fail to reject H → Median residual = 0 → No central bias.
Median residual: -0.000374
Wilcoxon statistic: 60620.0000
p-value: 0.933142
Fail to reject H → Median residual = 0 → No central bias.
Median residual: -0.962427
Wilcoxon statistic: 14906.0000
p-value: 0.000000
Reject H → Median residual 0 → Some central bias.

```

### 1.2.27 Interpretation – Wilcoxon Signed-Rank Test (Median Residual Bias)

The Wilcoxon test checks whether the **median residual = 0**, meaning no central (median) bias.

#### Model 1

- Median residual 0.38
- p-value = 0.059 → **Fail to reject H**  
→ No significant median bias.

#### Model 2

- Median residual -2.88
- p-value = 0.823 → **Fail to reject H**  
→ No central bias despite the negative mean.

#### Model 3

- Median residual 0
- p-value = 0.933 → **Fail to reject H**  
→ Residuals are centered around zero; no median bias.

#### Model 4

- Median residual -0.96
  - p-value = 0.000 → **Reject H**  
→ Median residual 0 → **Model shows significant central bias**  
→ Predictions systematically deviate from the true center.
-

### 1.2.28 Overall Summary

- **No median bias:** Models 1, 2, 3

- **Significant median bias:** Model 4

This shows only the last model consistently shifts predictions away from the true central value.

[ ]:

[ ]:

[ ]:

[ ]:

[ ]:

[ ]: